



¹Georgia Tech



²Univ of Toronto



³Texas A&M



⁴Univ of Wisconsin



SEEC: Stochastic Escape Express Channel

Mayank Parasar¹, Natalie Enright Jerger², Paul V. Gratz³,
Joshua San Miguel⁴, Tushar Krishna¹



<https://github.com/noc-deadlock/seec>

Presented By:
Tushar Krishna

Associate Professor

School of ECE

Georgia Institute of Technology

Email: tushar@ece.gatech.edu

Interconnection Networks are at the heart of HPC

News

Chinese 260-core processors ShenWei SW26010 enabled supercomputer Sunway TaihuLight be the most productive in the world

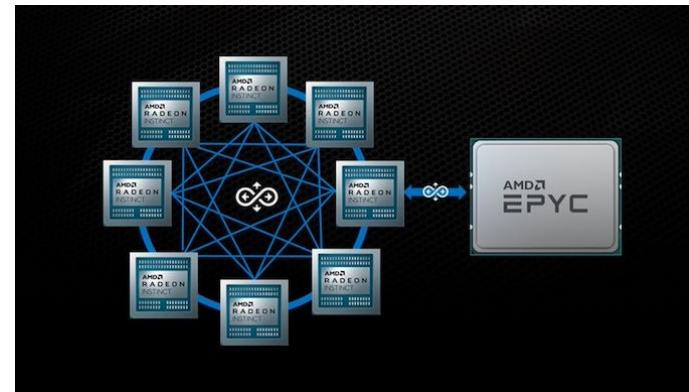
June 21, 2016 · 331 · 0



AMD's Third-Gen Infinity Architecture Enables Coherent CPU-GPU Communication

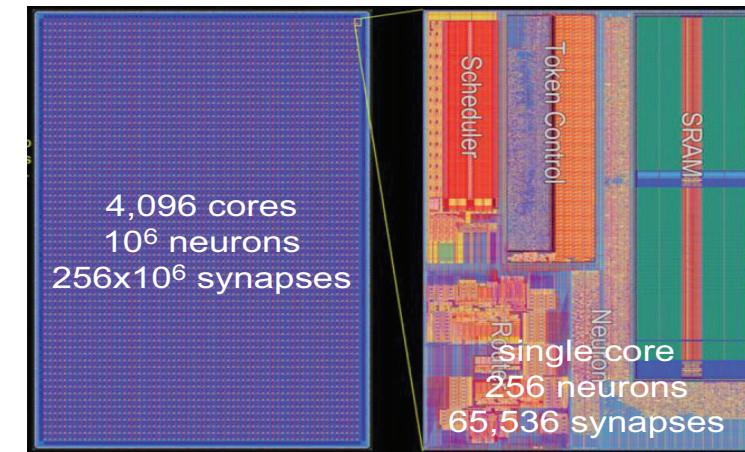
By Francisco Pires 8 days ago

The connecting tissue for Multi-Chip-Modules and Epyc HPC



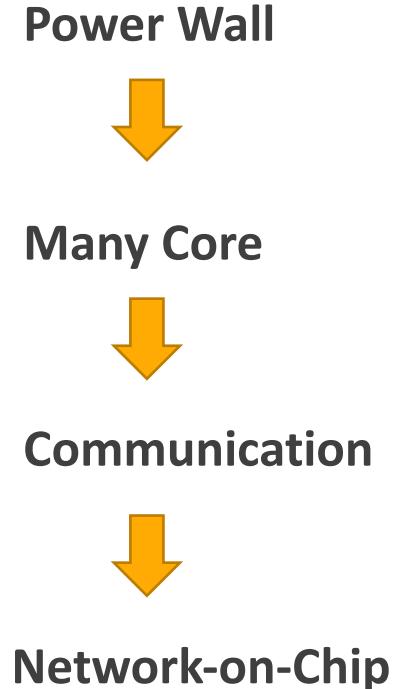
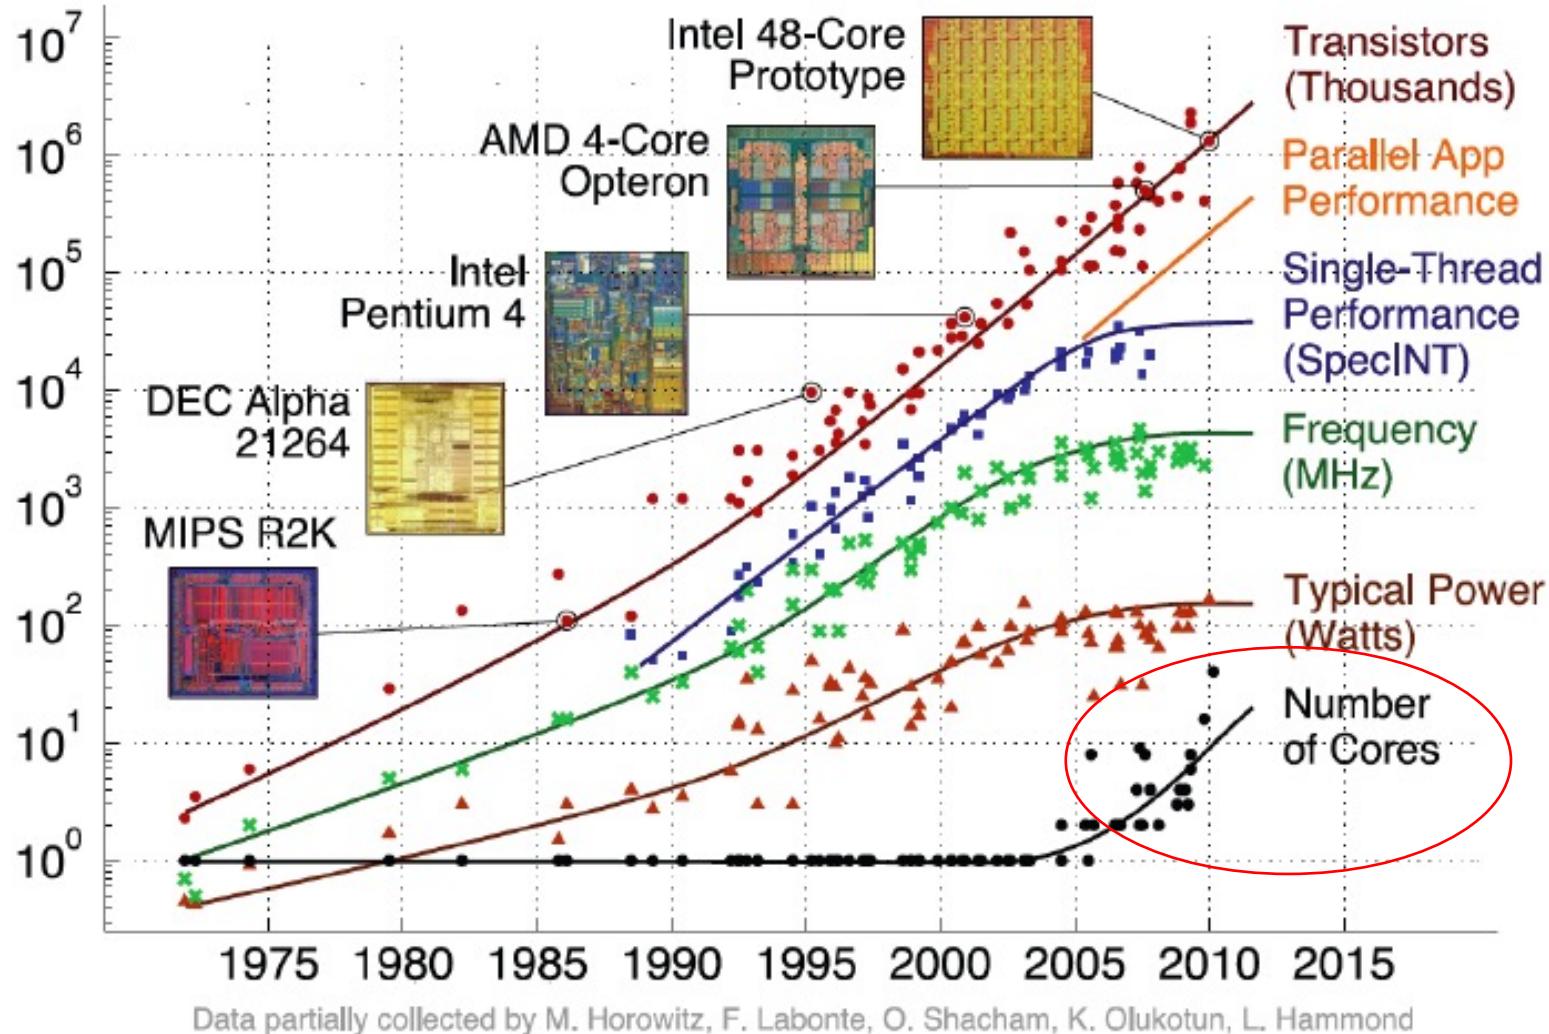
Technology

IBM reveals 'brain-like' chip with 4,096 cores



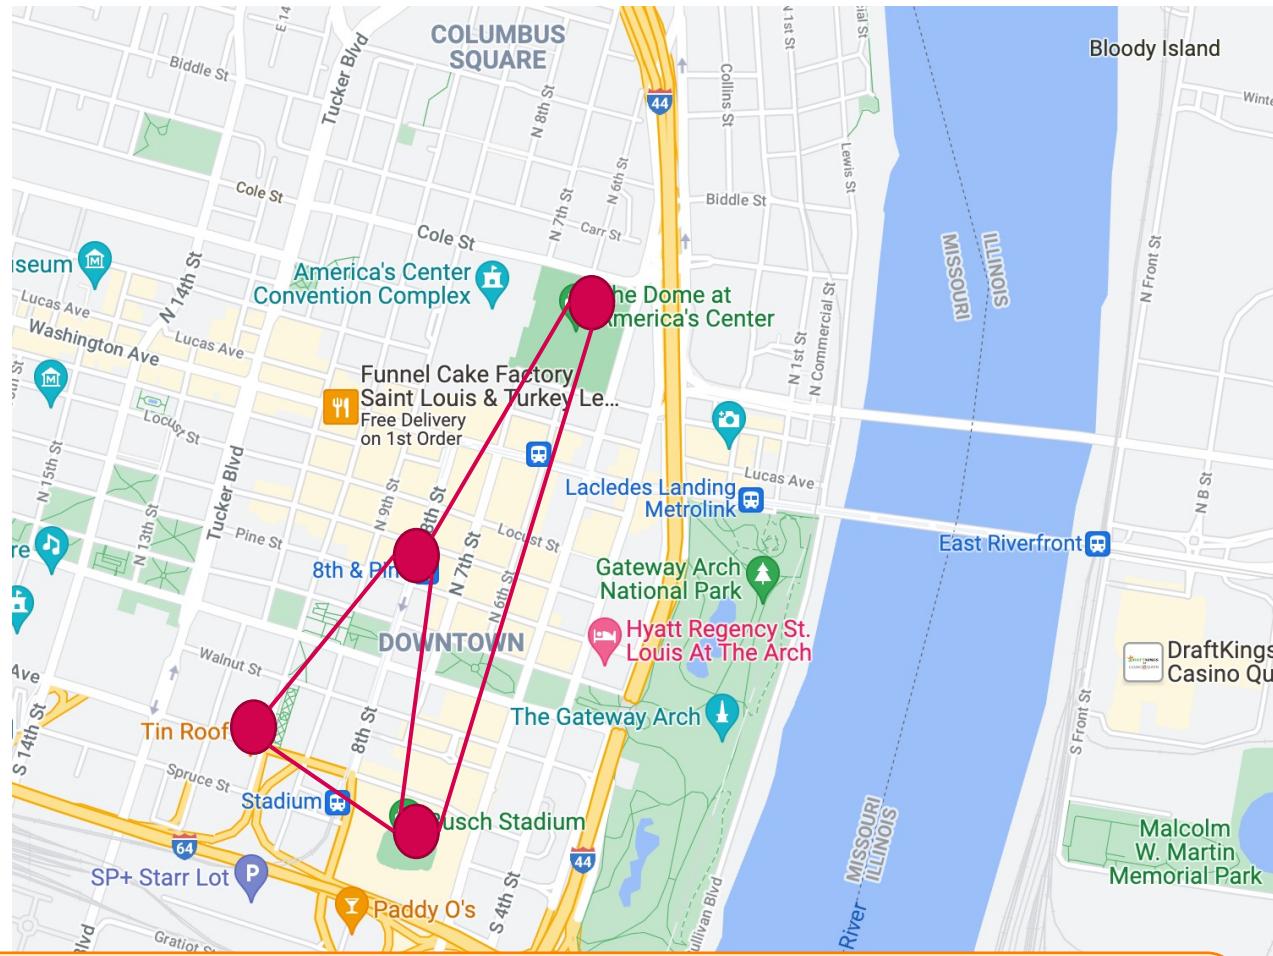
This work focuses on interconnection networks-on-chip (aka NoCs)

Why do we need Networks-on-Chip?



Network-on-Chip 101

- **Topology**
 - How to connect the nodes
 - *~Road Network*
- **Routing**
 - Which path should a message take
 - *~Series of road segments from source to destination*
- **Flow Control**
 - When does the message have to stop/proceed
 - *~Traffic signals at end of each road segment*
- **Router Microarchitecture**
 - How to build the routers
 - *~Design of traffic intersection (number of lanes, algorithm for turning red/green)*



Design-time **hardware-managed** topology, routing and flow-control decisions makes networks-on-chip different from larger scale HPC / datacenter networks

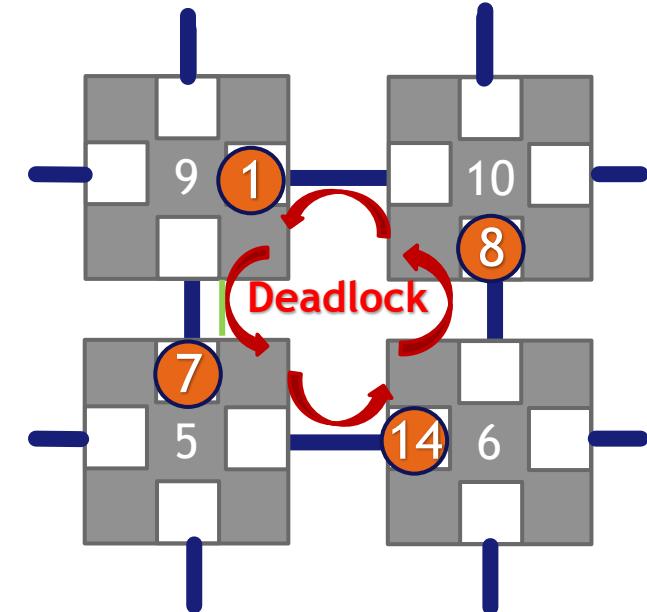
Challenges with NoC Design

- Correctness → Guaranteed packet delivery

- No Routing Deadlocks
- No Protocol Deadlocks



Challenge: Cyclic dependence



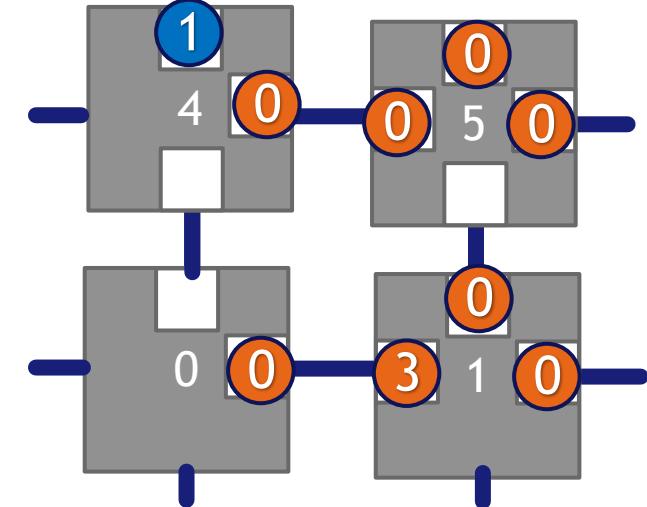
- Performance → Deliver as many packets as possible

- Reduce Latency
- Increase Throughput



Challenge: Congestion

This work addresses both challenges with a unified solution



Outline

- **Background**

- Networks-on-Chip
- Deadlock-Freedom Techniques

- **SEEC**

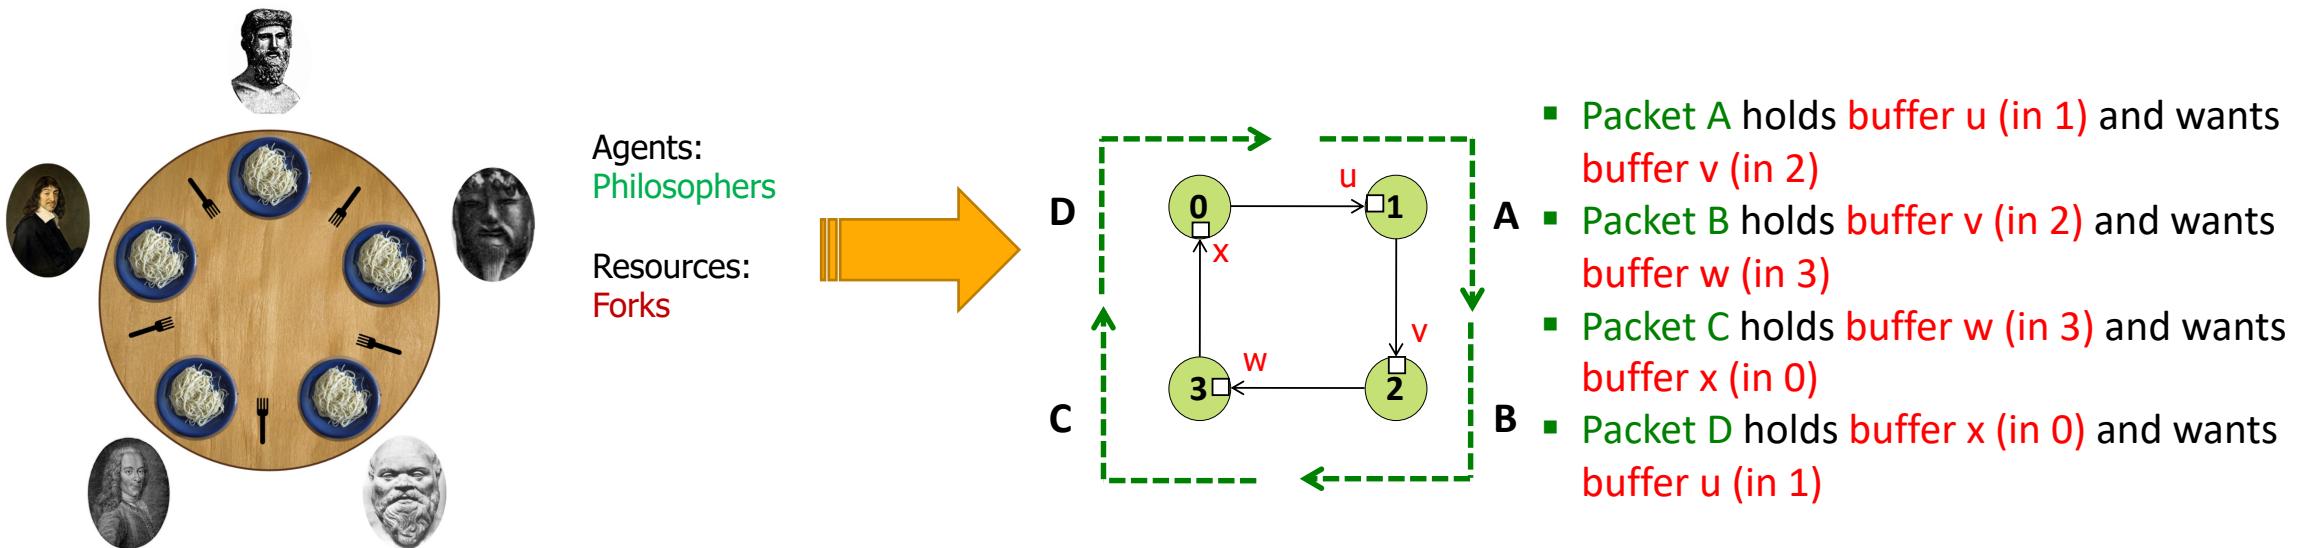
- Overview
- Walkthrough
- Features
- Optimizations

- **Evaluations**

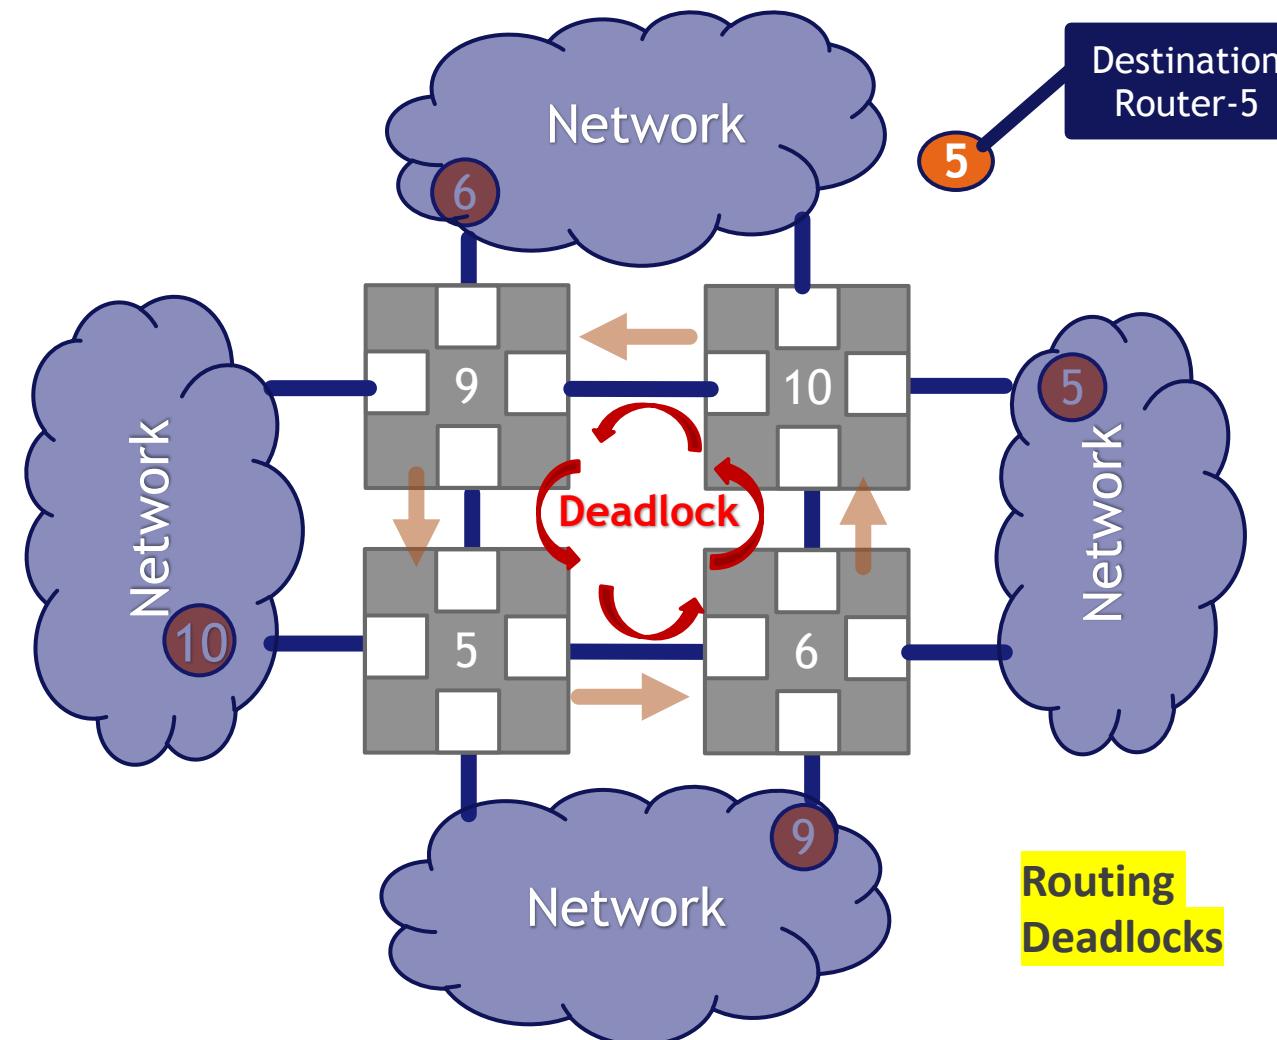
- **Conclusions**

What are deadlocks?

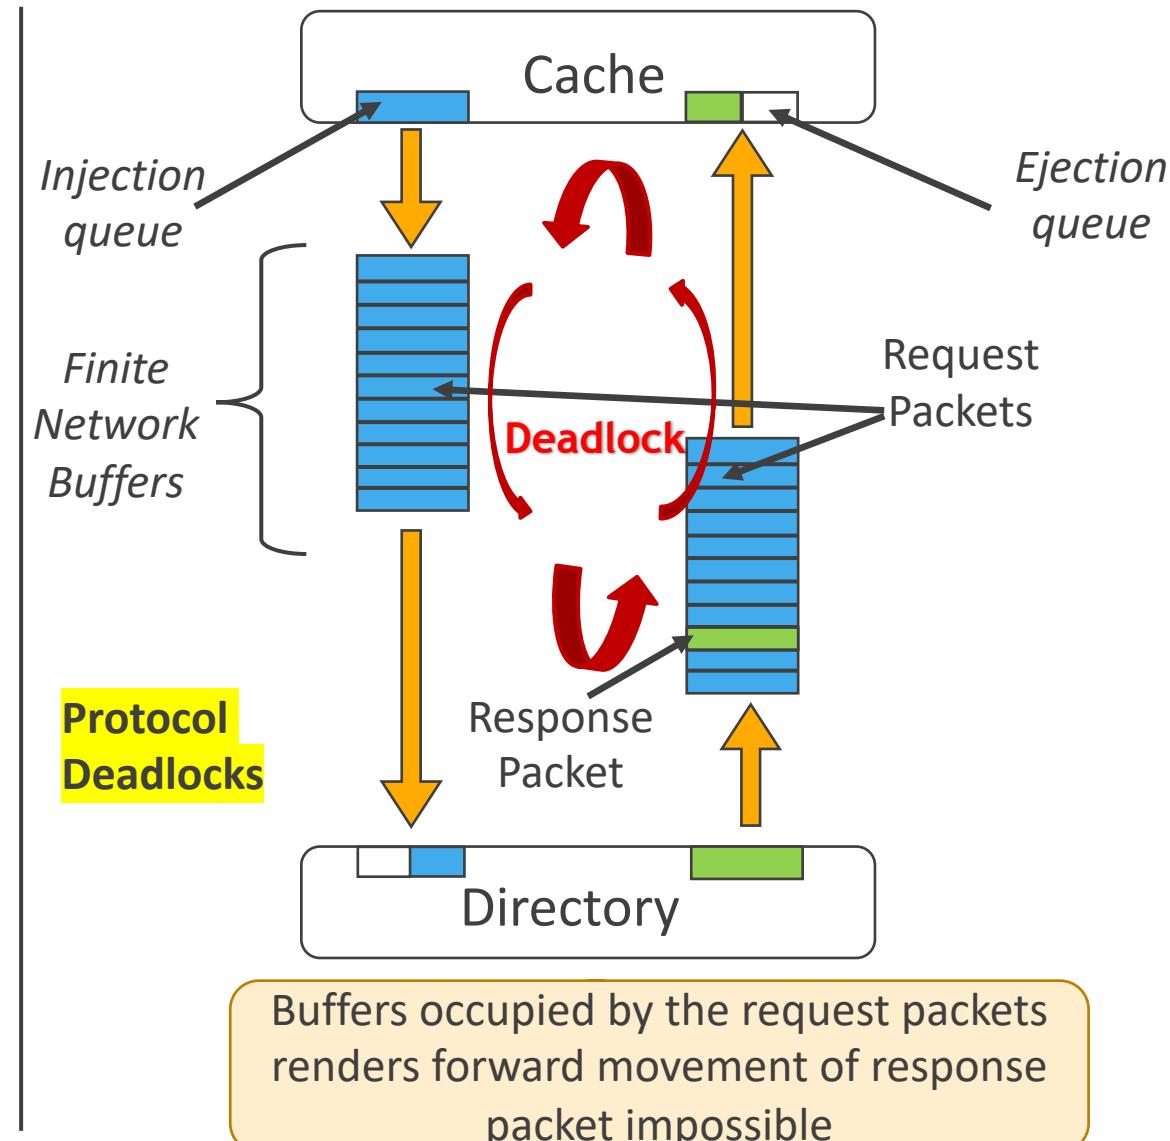
A condition in which a set of **agents** wait indefinitely trying to acquire a set of **resources**



Two kinds of deadlocks



**Routing
Deadlocks**



Current solutions for deadlock-freedom

- **Proactive Techniques**

- Avoid deadlock to begin with



Add turn restrictions or provide additional virtual lanes/channels

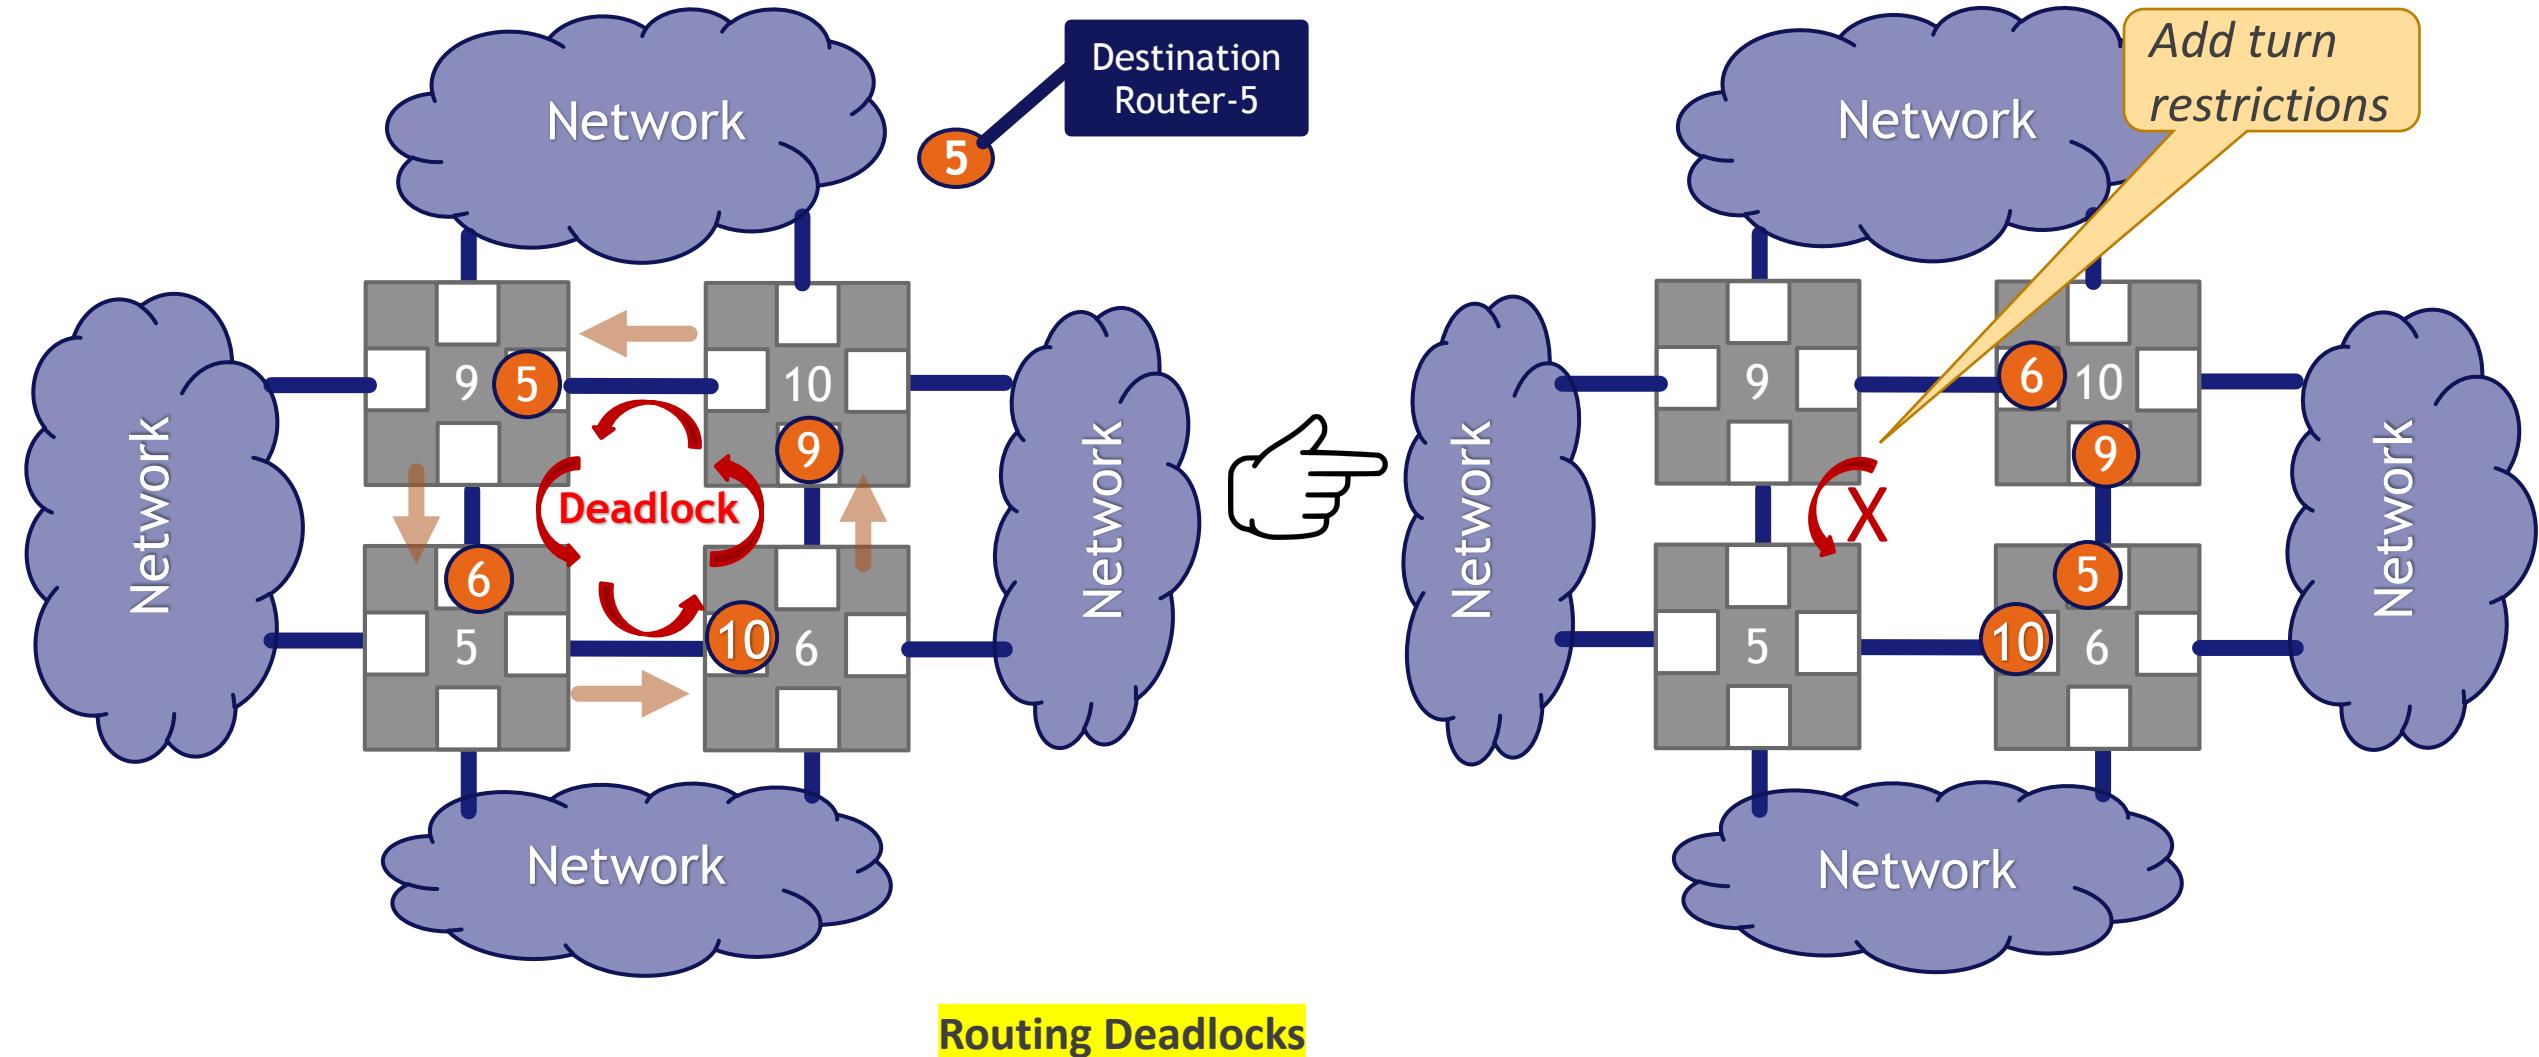
- **Reactive Techniques**

- Allow deadlocks to form in the network
 - Detect deadlocks and recover from them

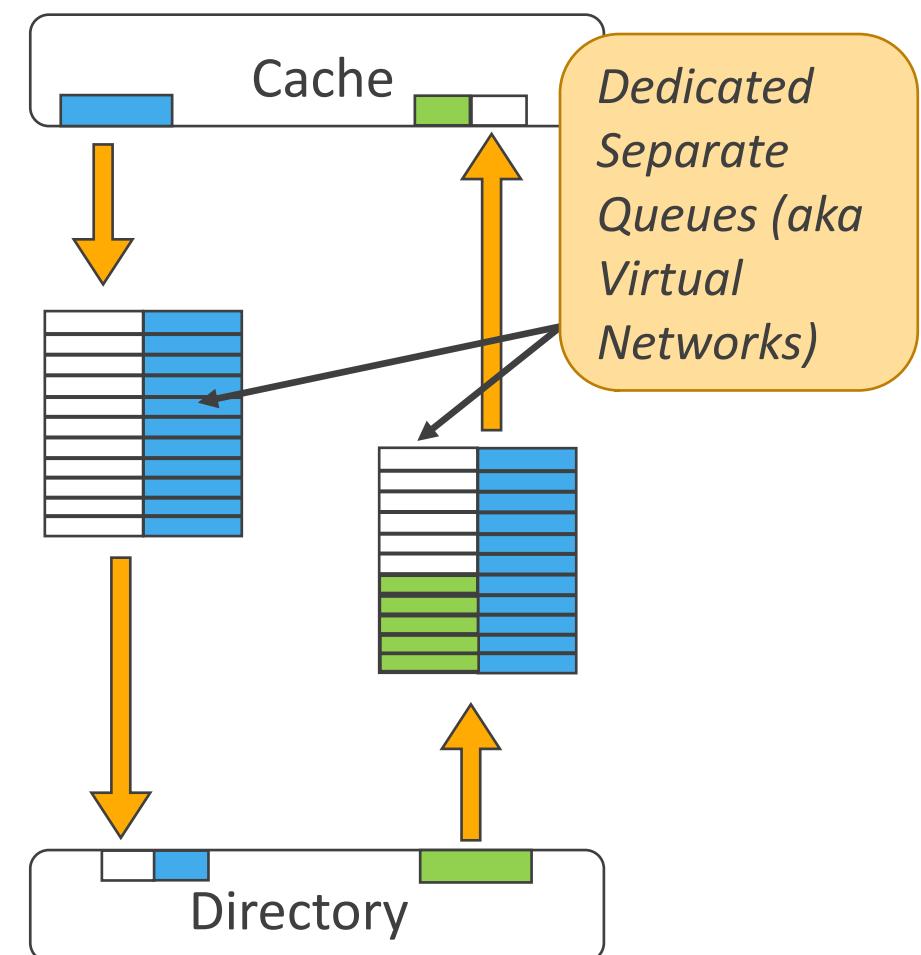
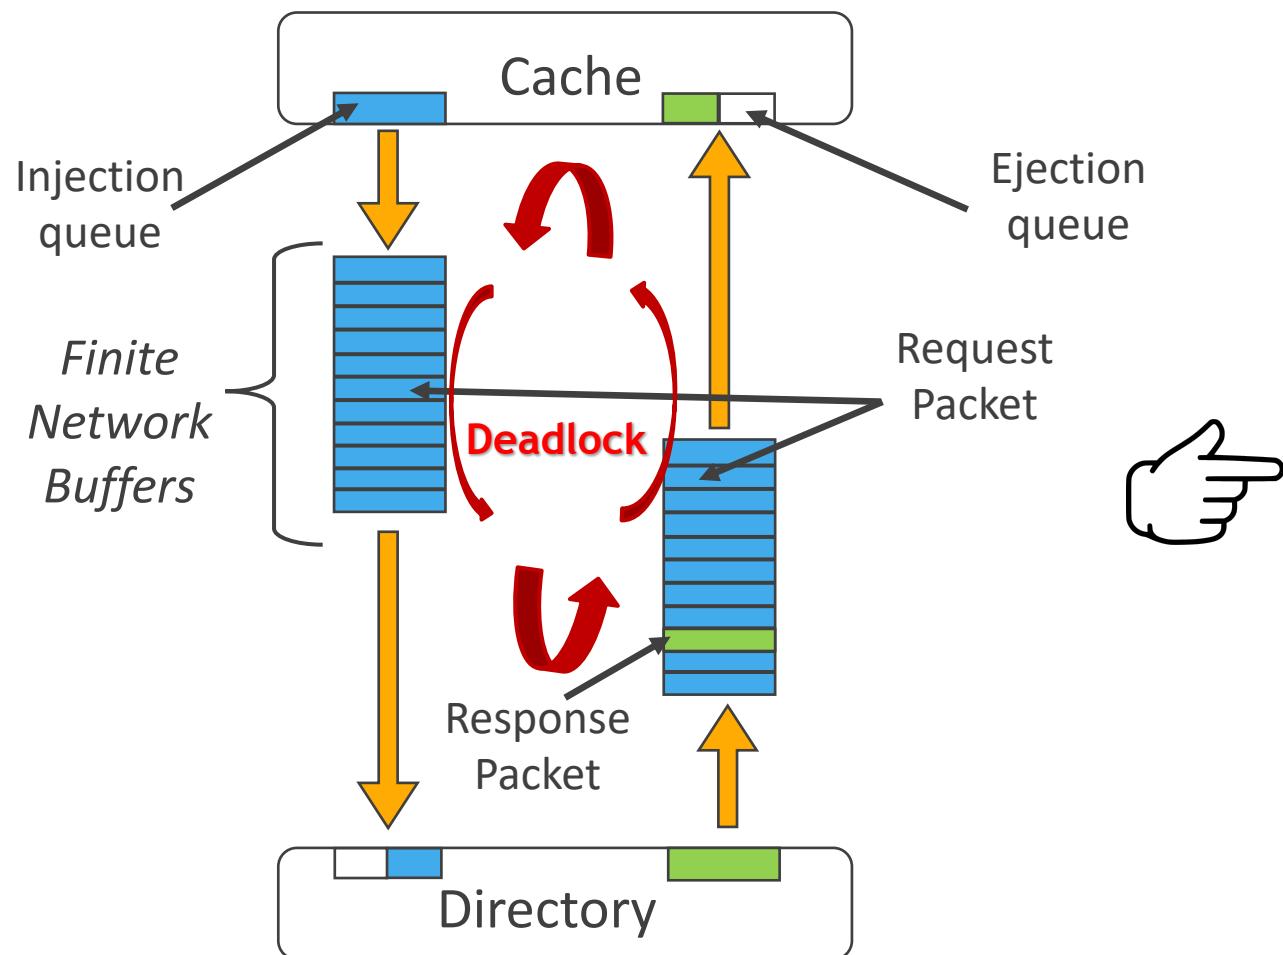


Add circuitry for monitoring and identifying cyclic dependencies at runtime

Example of proactive routing deadlock-freedom



Example of proactive protocol deadlock-freedom



Protocol Deadlocks

Deadlock-Freedom Techniques

Technique	Type of Solution	High Throughput	Low Area Power	Low Hardware Complexity	Routing Deadlock-free	Protocol Deadlock-free
Turn Restriction [1]	Proactive	✗	✓	✓	✓	✗
Escape VC [2]	Proactive	✓	✗	✓	✓	✗
Virtual Network [3]	Proactive	✓	✗	✓	✗	✓
Deadlock Recovery [4]	Reactive	✓	✓	✗	✓	✓

Can we do better?

- [1]: Deadlock-free message routing in multiprocessor interconnection networks
[IEEE Trans. Computer, 1987]
- [2]: A new theory of deadlock-free adaptive routing in wormhole networks.
[IEEE Trans. On Parallel and Dist. Sys, 1993]
- [3]: The Alpha 21364 network architecture [MICRO 2002]
- [4]: DISHA [IPDPS'95] Static Bubble [HPCA'17] SPIN [ISCA'18]

Takeaway:

- Proactive techniques add performance and area overheads.
- Reactive techniques require expensive monitoring circuitry

“Subactive” Techniques

- **Proactive Techniques**

- Avoid deadlock to begin with

- **Reactive Techniques**

- Allow deadlock to form in the network
 - Detect deadlocks and recover from them

- **Subactive Techniques** 

- Allow deadlocks to form in the network
 - Obliviously flush network to remove any deadlock

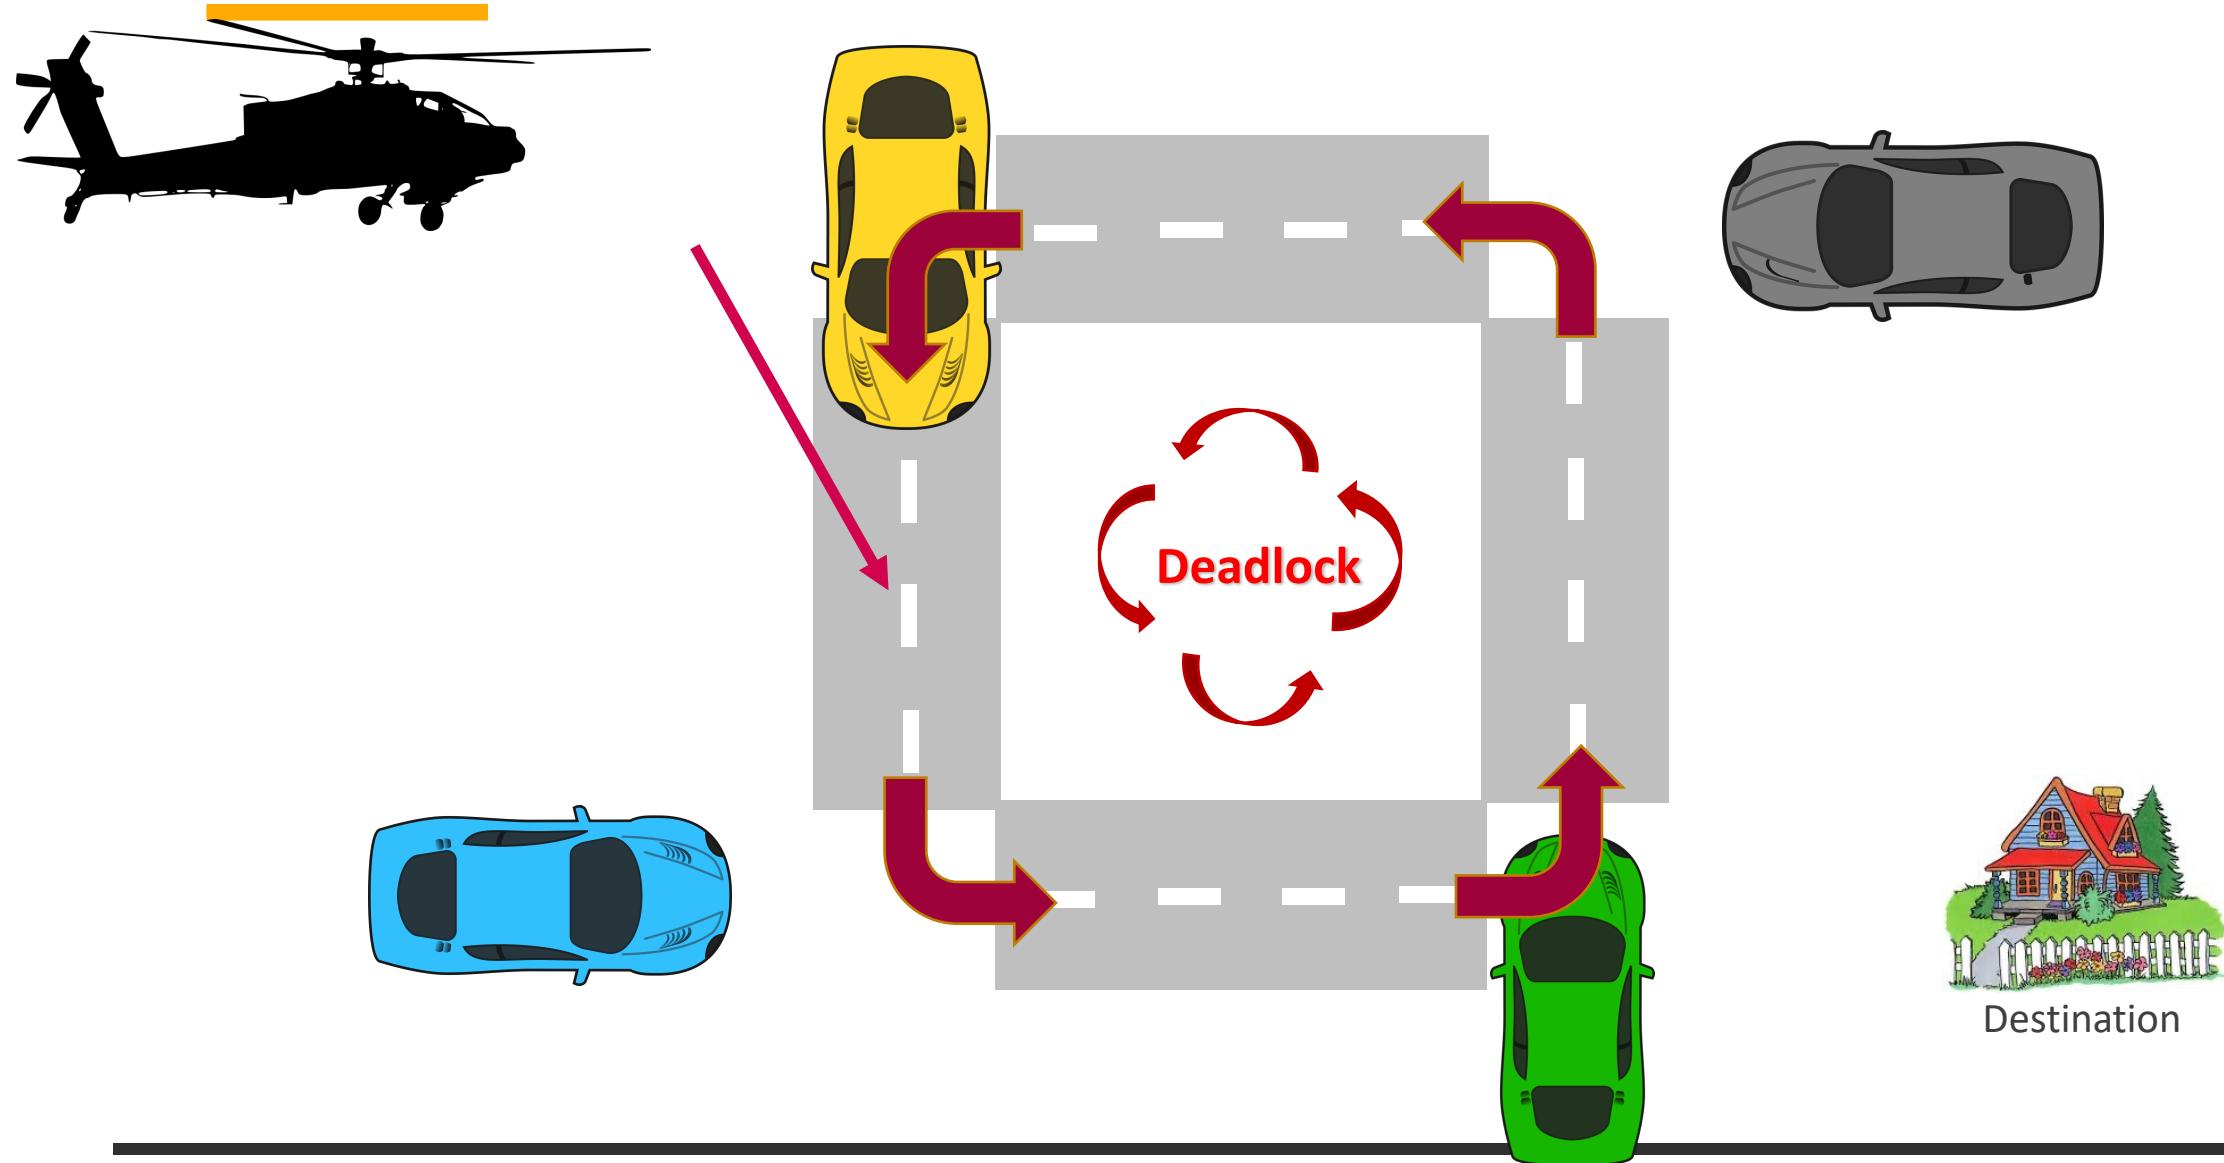
Intuition behind Subactive

- Deadlocks are rare!
 - Requires certain ensemble of packets with a particular buffer dependency for the deadlock to occur
- A mechanism to ensure periodic draining of packets **sufficient** to guarantee deadlock-freedom
 - No need for proactive performance restrictions / additional reserved queues
 - No need for detection and reactive recovery
- This work (SEEC) is one of many possible subactive solutions
 - See our other work: SWAP (MICRO 2019), DRAIN (HPCA 2020), PitStop (HPCA 2021)

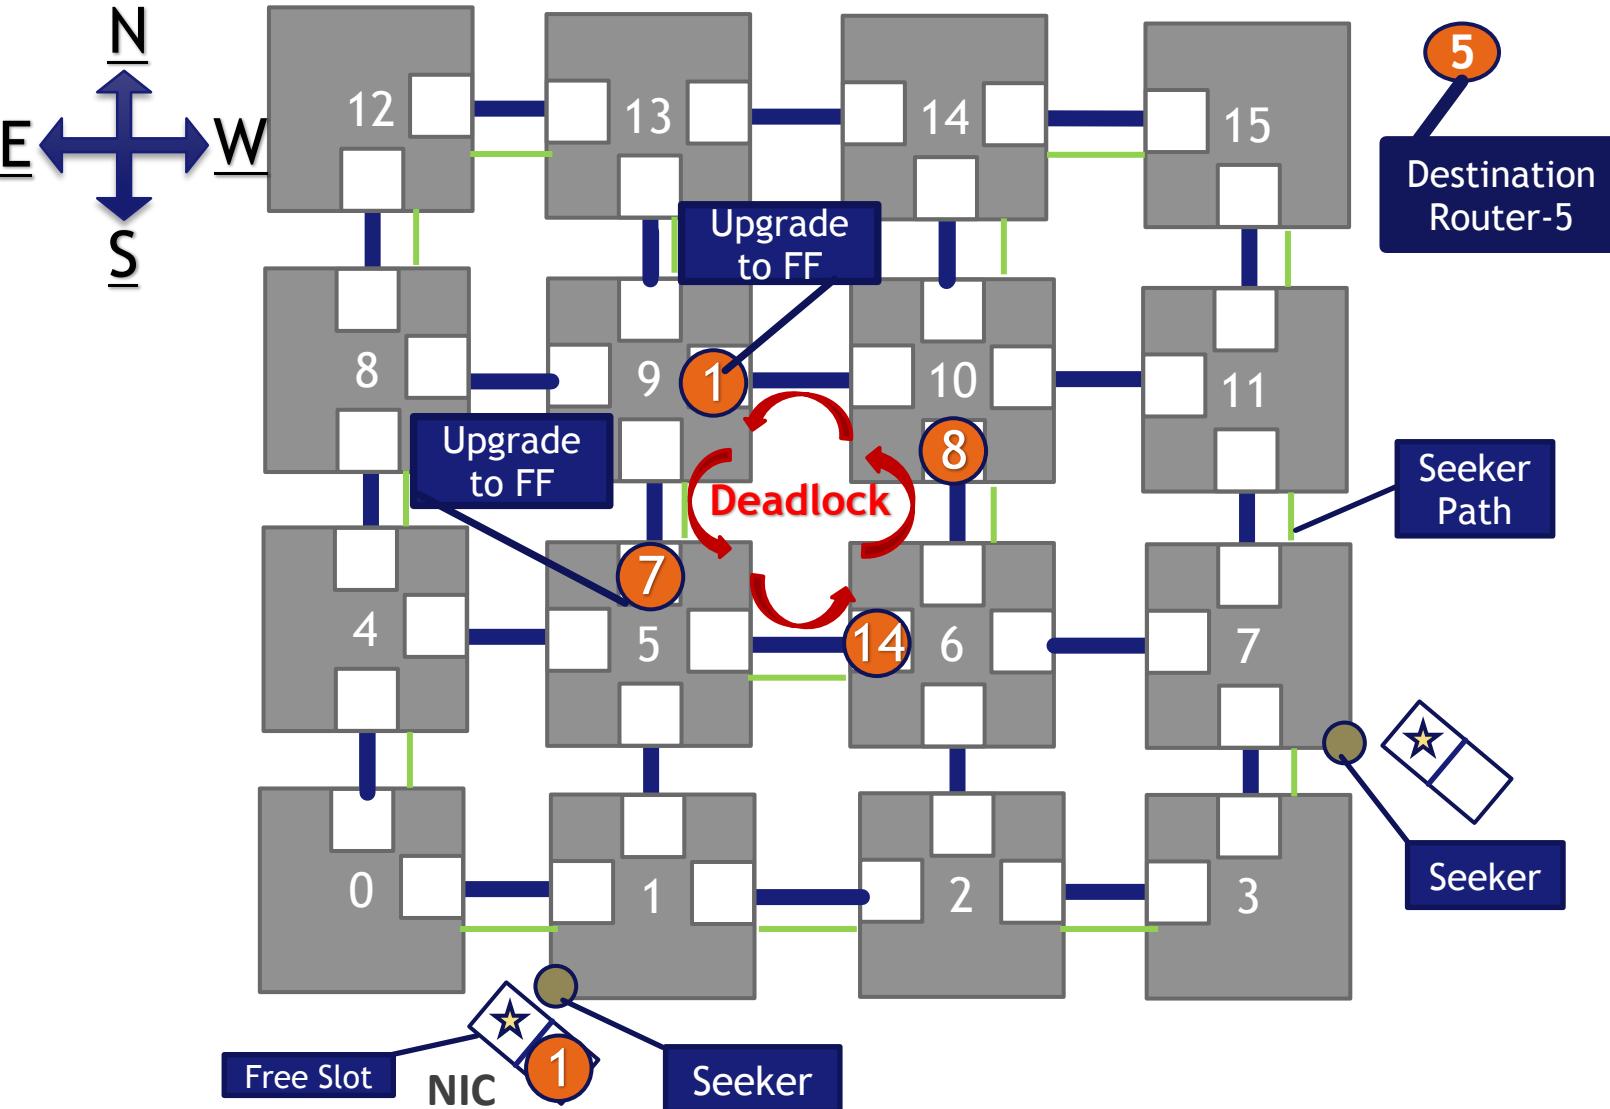
Outline

- **Background**
 - Networks-on-Chip
 - Deadlock-Freedom Techniques
- **SEEC**
 - Overview
 - Walkthrough
 - Features
 - Optimizations
- **Evaluations**
- **Conclusions**

SEEC: Deadlock Analogy



SEEC: Walk-through for Deadlock-freedom



1. Reserve a free ejection buffer at the NIC
2. Send a *seeker* over a sideband *seeker path* through all routers
3. If a packet intended for this NIC is found, upgrade it for ***Free Flow (FF)***
4. FF packet given priority over regular packets to traverse all the way to the N
5. Repeat for all NICs in round-robin manner

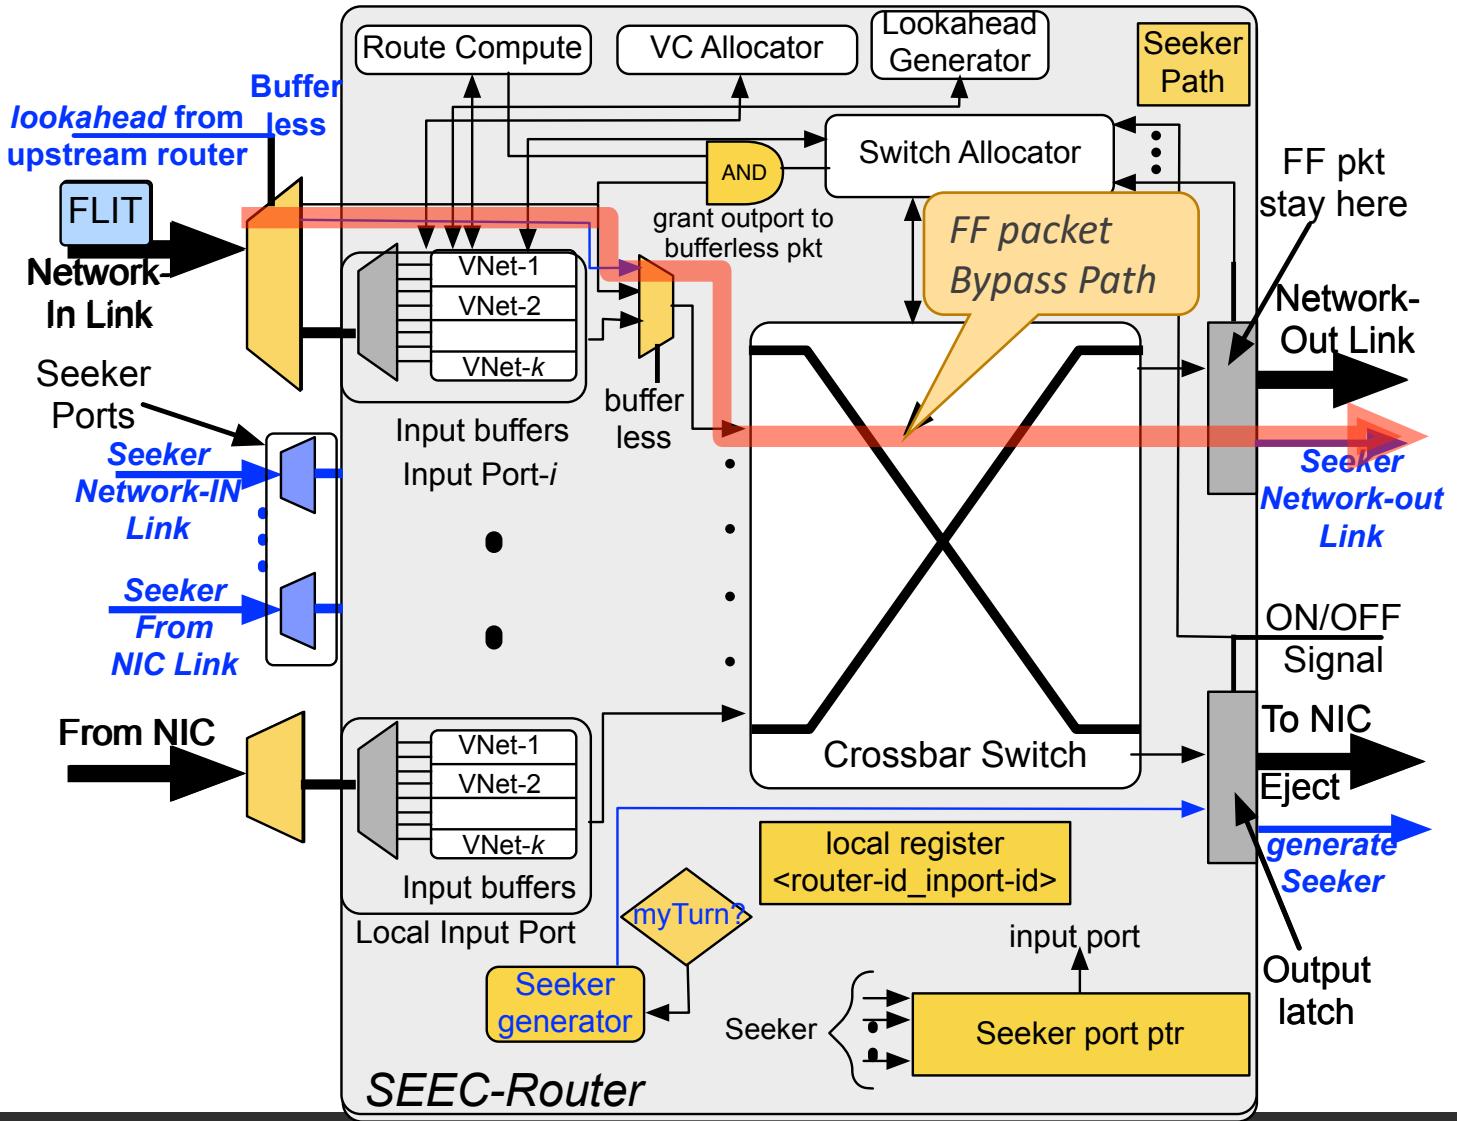
Outline

- **Background**
 - Networks-on-Chip
 - Deadlock-Freedom Techniques
- **SEEC**
 - Overview
 - Walkthrough
 - **Features**
 - Optimizations
- **Evaluations**
- **Conclusions**

Features of SEEC

- **Novel Flow Control: Free Flow (FF)**
 - Packets upgraded to FF via seeker
 - FF packet gets prioritized over the buffered packets
 - FF packet sends a *lookahead* signal on cycle in advance
 - *lookahead* signal sets the appropriate mux at downstream router
 - FF packet traverse the network from link to link till the NIC
 - FF packets follow ***minimal (shortest) express path*** with no buffering
 - SEEC proposes co-existence of FF packets and buffered packets
- **Seeker Path**
 - Any pre-defined path that visits all routers
 - Separate sideband network
 - could use regular links by prioritizing seeker over regular packets
 - Possible to have multiple parallel seekers → later in talk

Router Microarchitecture



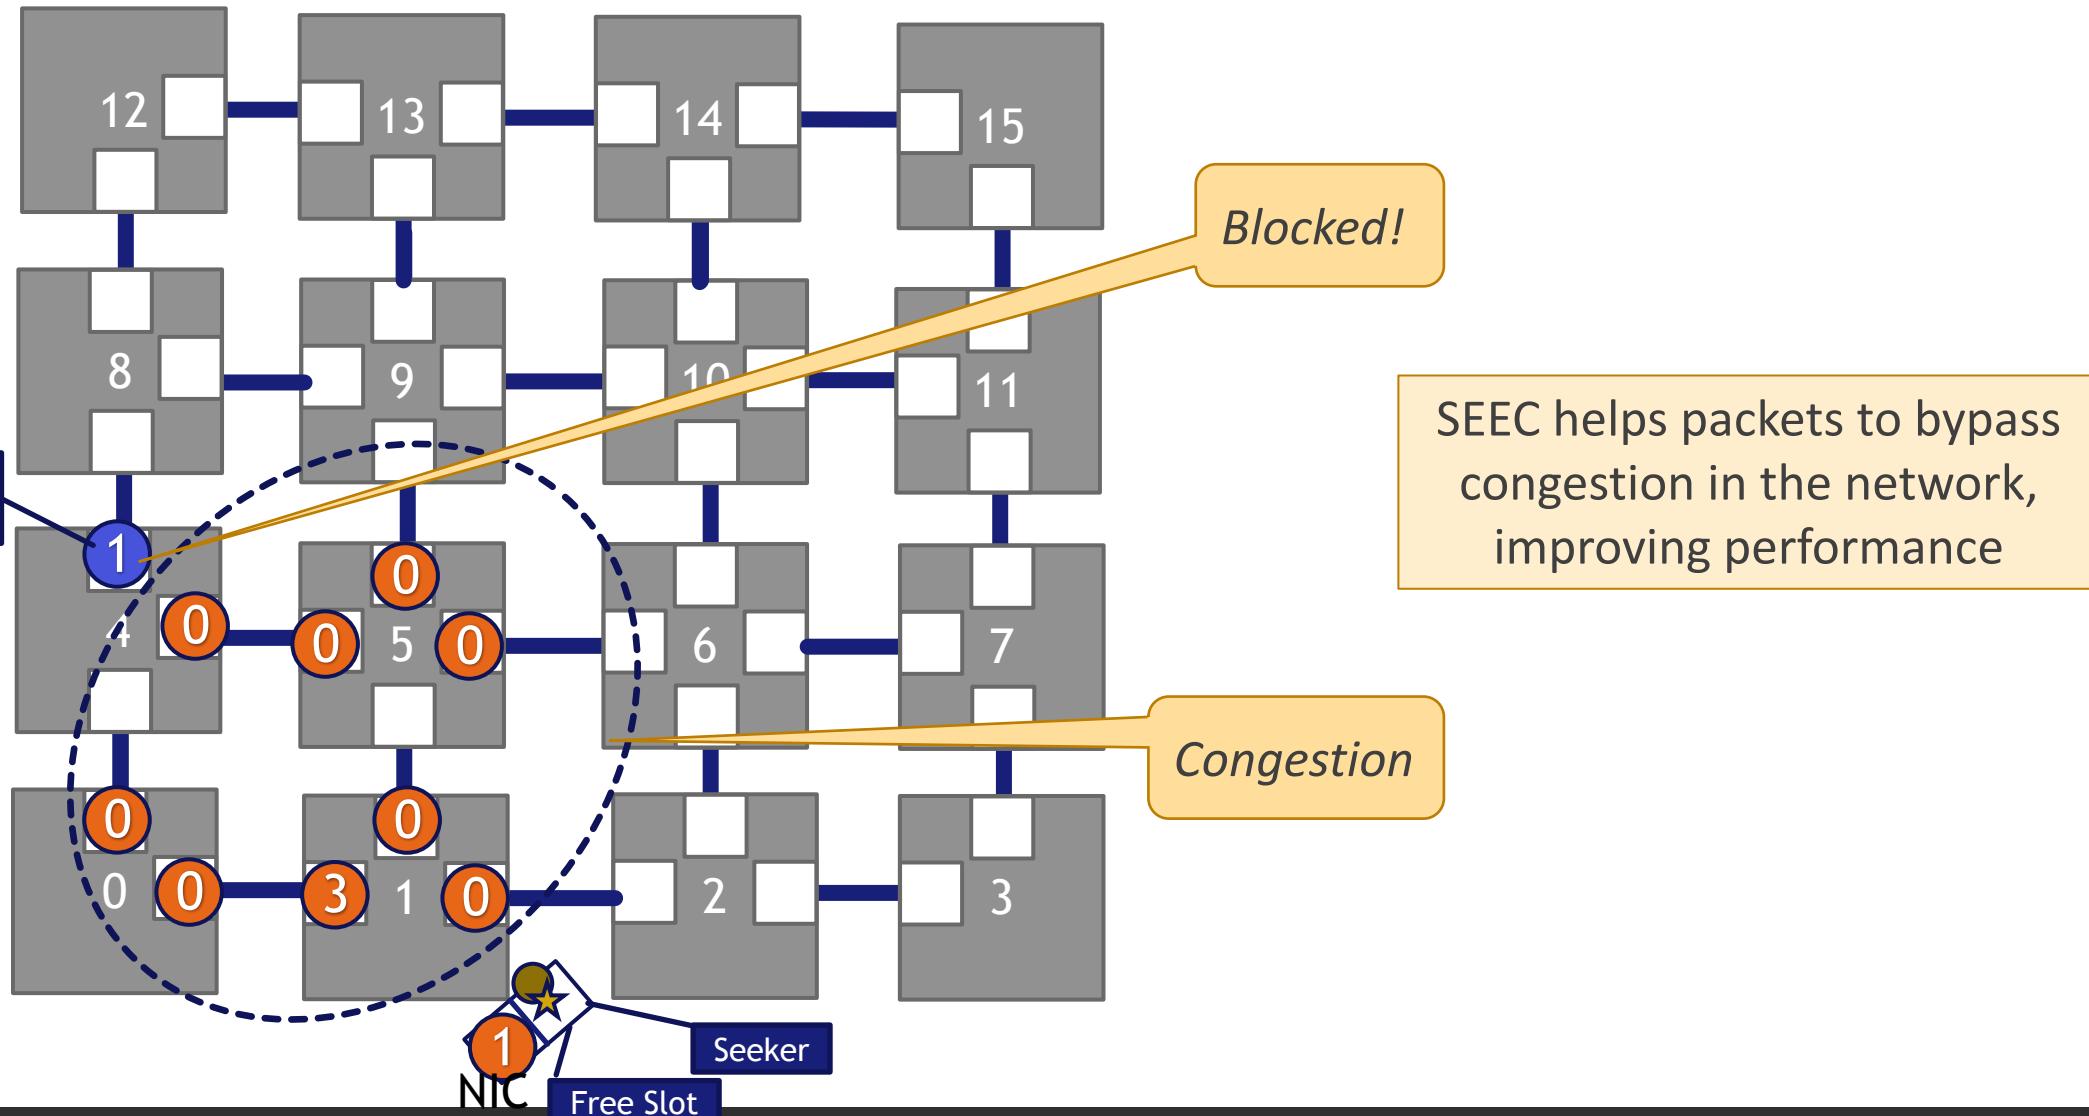
Deadlock Freedom Proof

- The ejection port at NIC has separate Virtual Networks (one per protocol message class)
- The response packet can cause a protocol deadlock, if stuck in network
- SEEC is ***protocol deadlock free*** as NIC will eject any blocked response packet by sending seeker
- SEEC is ***routing deadlock free*** as all the input ports get chance to send packets via FF flow control
- SEEC is ***livelock free*** as there is no misrouting

Mathematical upper bound for deadlock persistence

- For any irregular network topology:
 - Assume diameter is D
 - Assume it takes on an average ‘ k ’ cycles at the ejection port
 - Number of nodes in the network is ‘ N ’
 - Let the average network router-radix of the router is ‘ r ’
- The maximum time it takes for a deadlock to persist
 - Deadlock involves at least 4 routers
 - $\{r*(N-4)*2*D*k\} + (D+k)$ cycles

SEEC: Additional Benefits → Performance



SEEC helps packets to bypass congestion in the network, improving performance

Outline

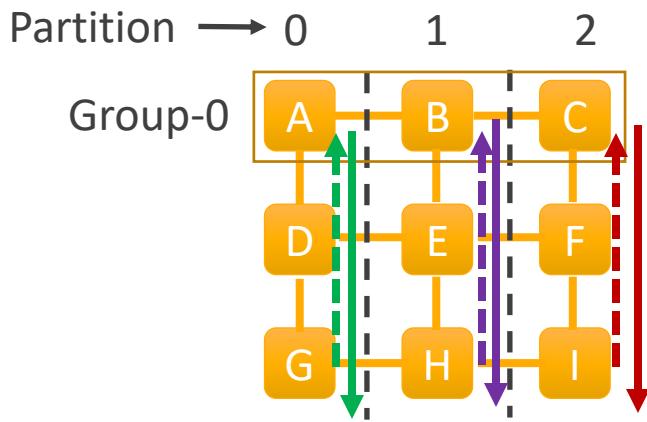
- **Background**
 - Networks-on-Chip
 - Deadlock-Freedom Techniques
- **SEEC**
 - Overview
 - Walkthrough
 - Features
 - Optimizations
- **Evaluations**
- **Conclusions**

Multi-SEEC (mSEEC)

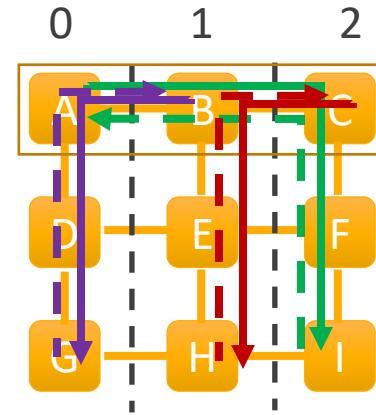
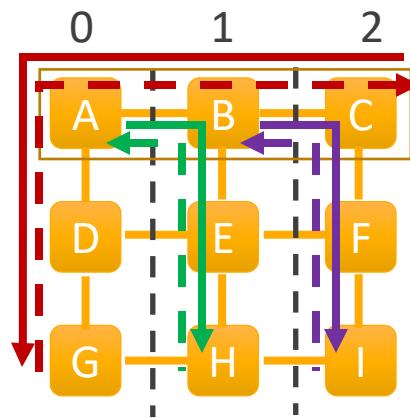
- **Motivation**
 - Path Diversity in many topologies: Multiple ways of reaching a node
- **Features of mSEEC**
 - Creates multiple FF packets that can traverse simultaneously
- **Requirement**
 - Multiple FF packets traverse minimally without overlapping paths
- **Implementation**
 - mSEEC divides the topology into independent regions

mSEEC: Walk Through

Phase 0



→ Seeker → Packet

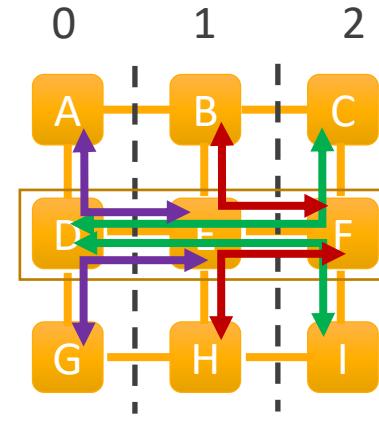
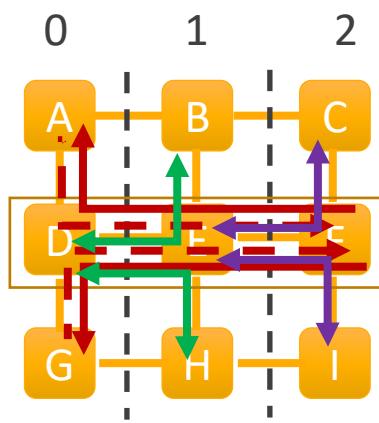
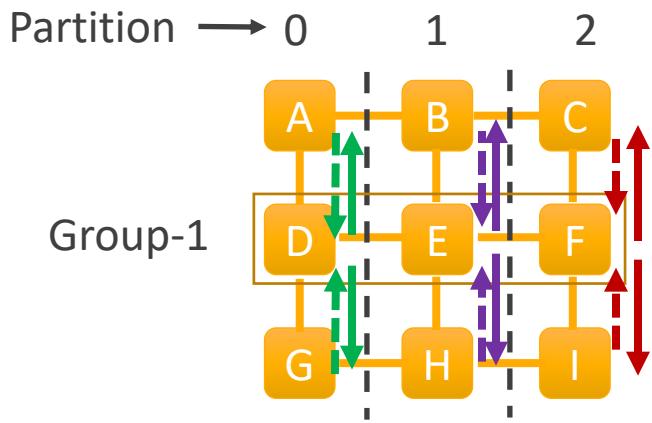


- When Phase-0 finishes:
 - each node in group-0 has searched entire topology
- Phase-1 starts

mSEEC: Walk Through

Phase 1

→ Seeker - - - → Packet



- When Phase-1 finishes:
 - each node in group-1 has searched entire topology
- Phase-2 starts and so on...

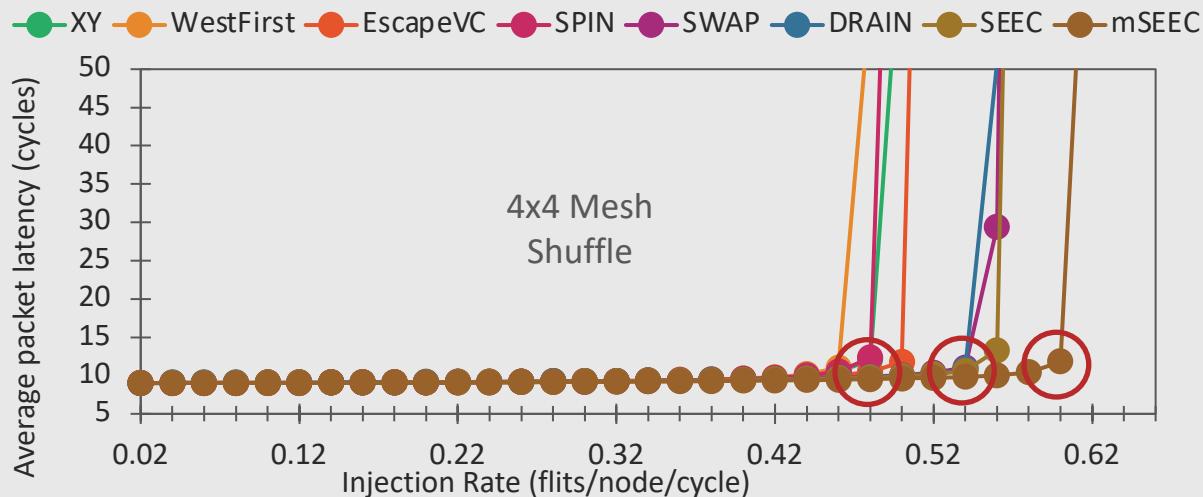
Outline

- **Background**
 - Networks-on-Chip
 - Deadlock-Freedom Techniques
- **SEEC**
 - Overview
 - Walkthrough
 - Features
 - Optimizations
- **Evaluations**
- Conclusions

Evaluations Methodology

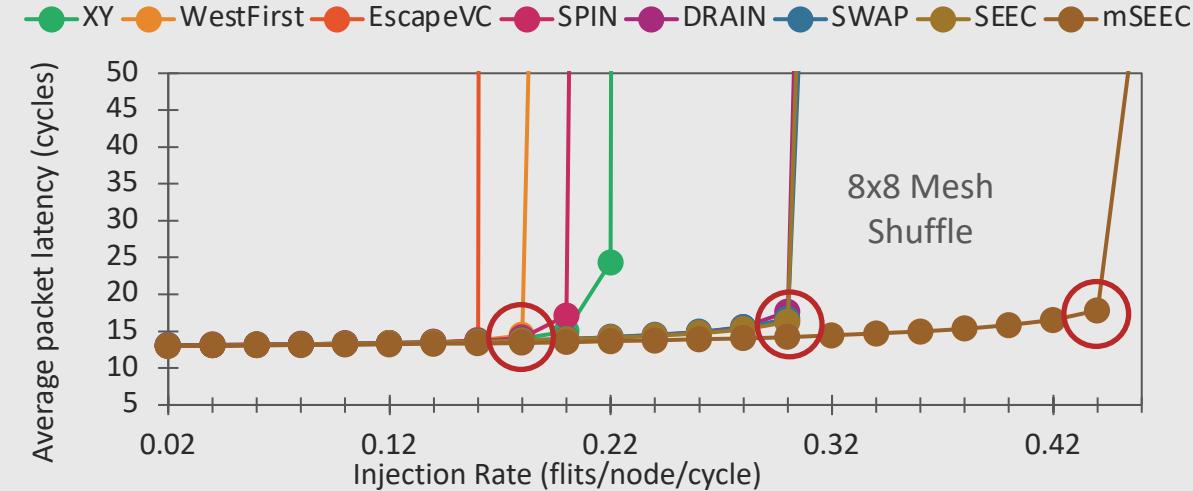
- gem5 for full system and synthetic simulations
- Garnet2.0 for network simulation
- SEEC/mSEEC uses minimal adaptive routing algorithm with full path diversity
- Virtual Cut Through Buffer Management
- Evaluated on 4x4 (16 nodes), 8x8 (64 nodes), and 16x16 (256 nodes) Mesh topology

Results: Performance Sweep



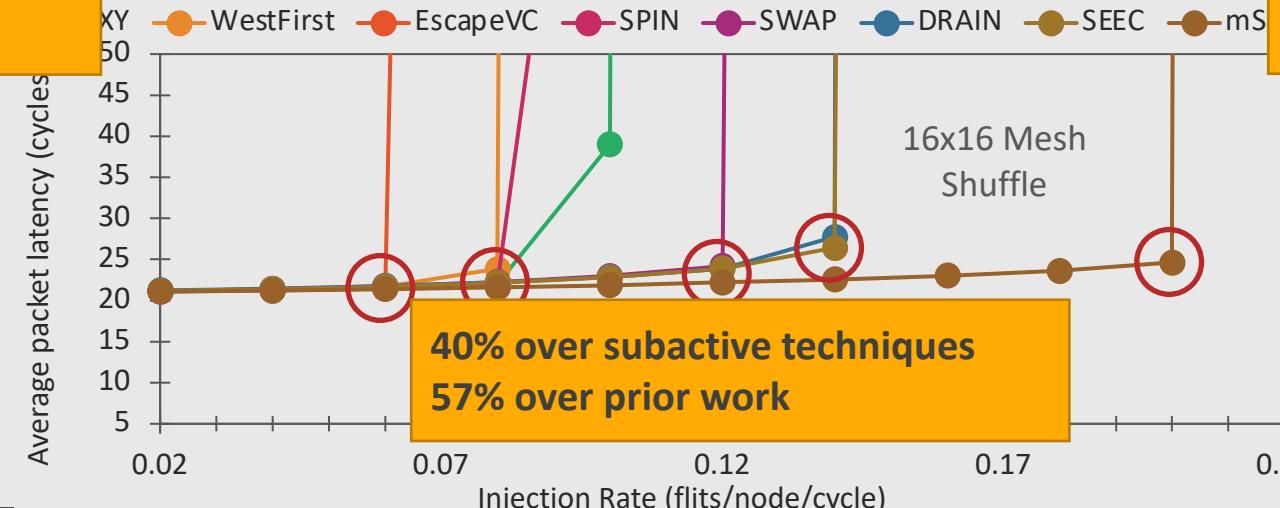
11% over subactive techniques

25% over prior work



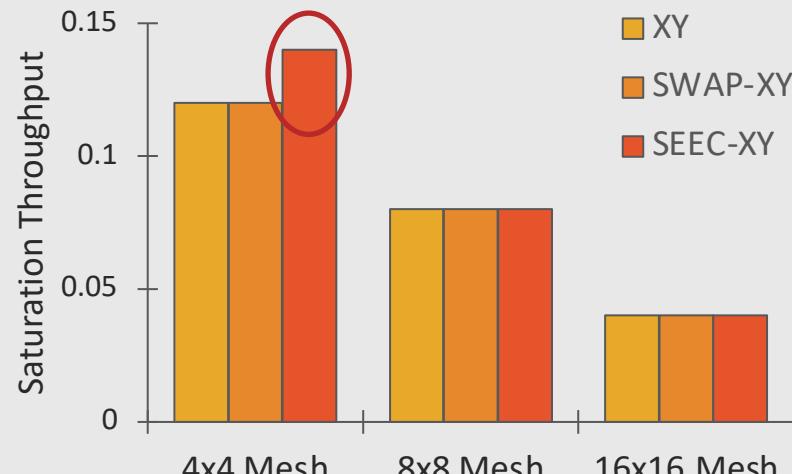
30% over subactive techniques

40% over prior work

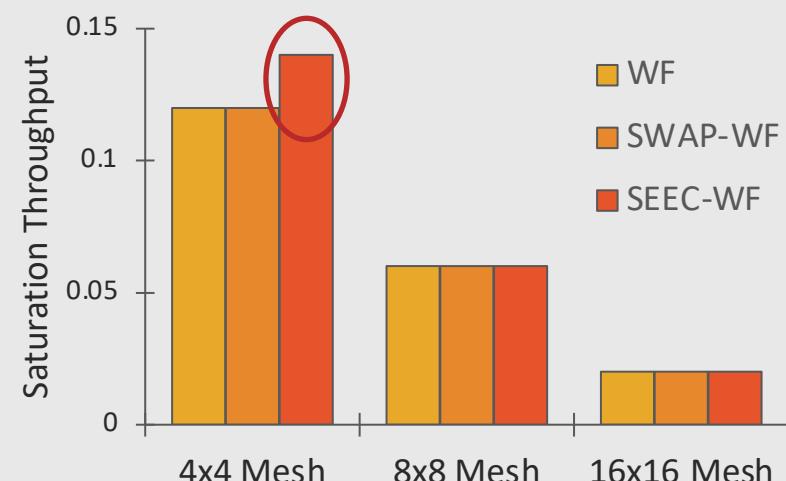


40% over subactive techniques
57% over prior work

Results: Throughput Enhancement



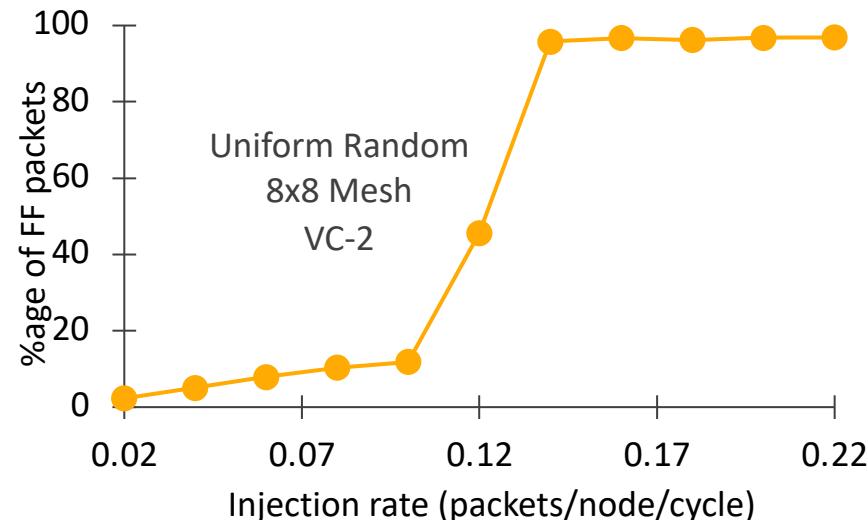
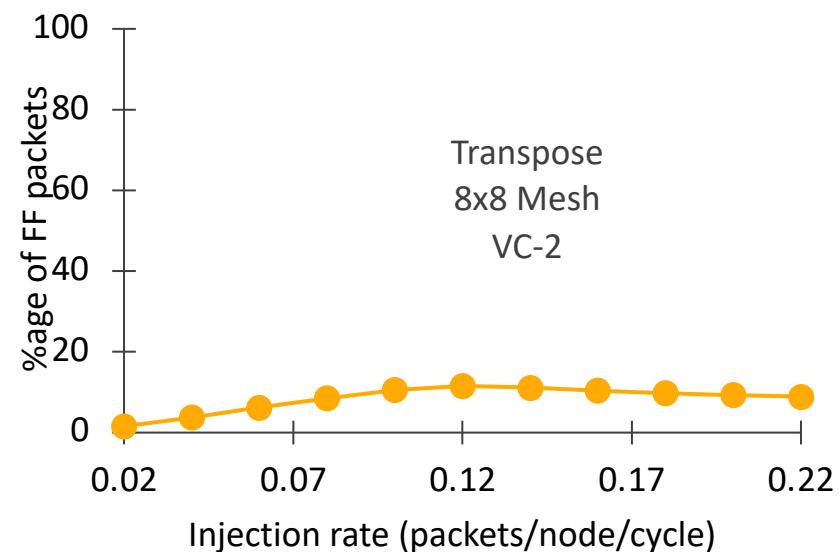
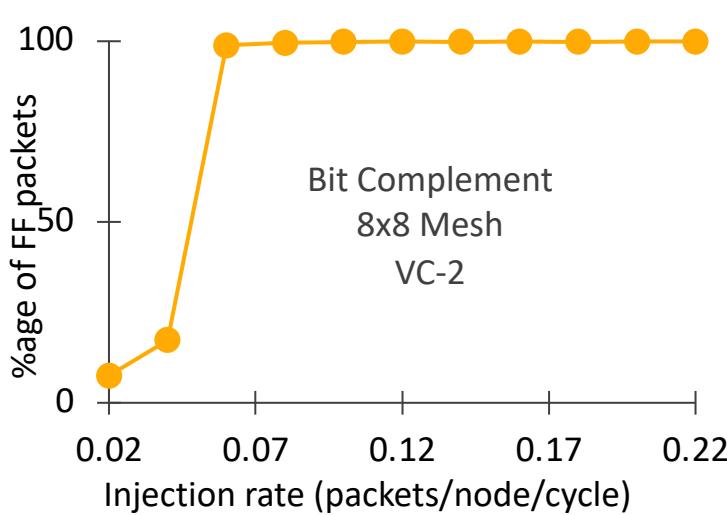
Uniform Random
XY Routing



Uniform Random
West First Routing

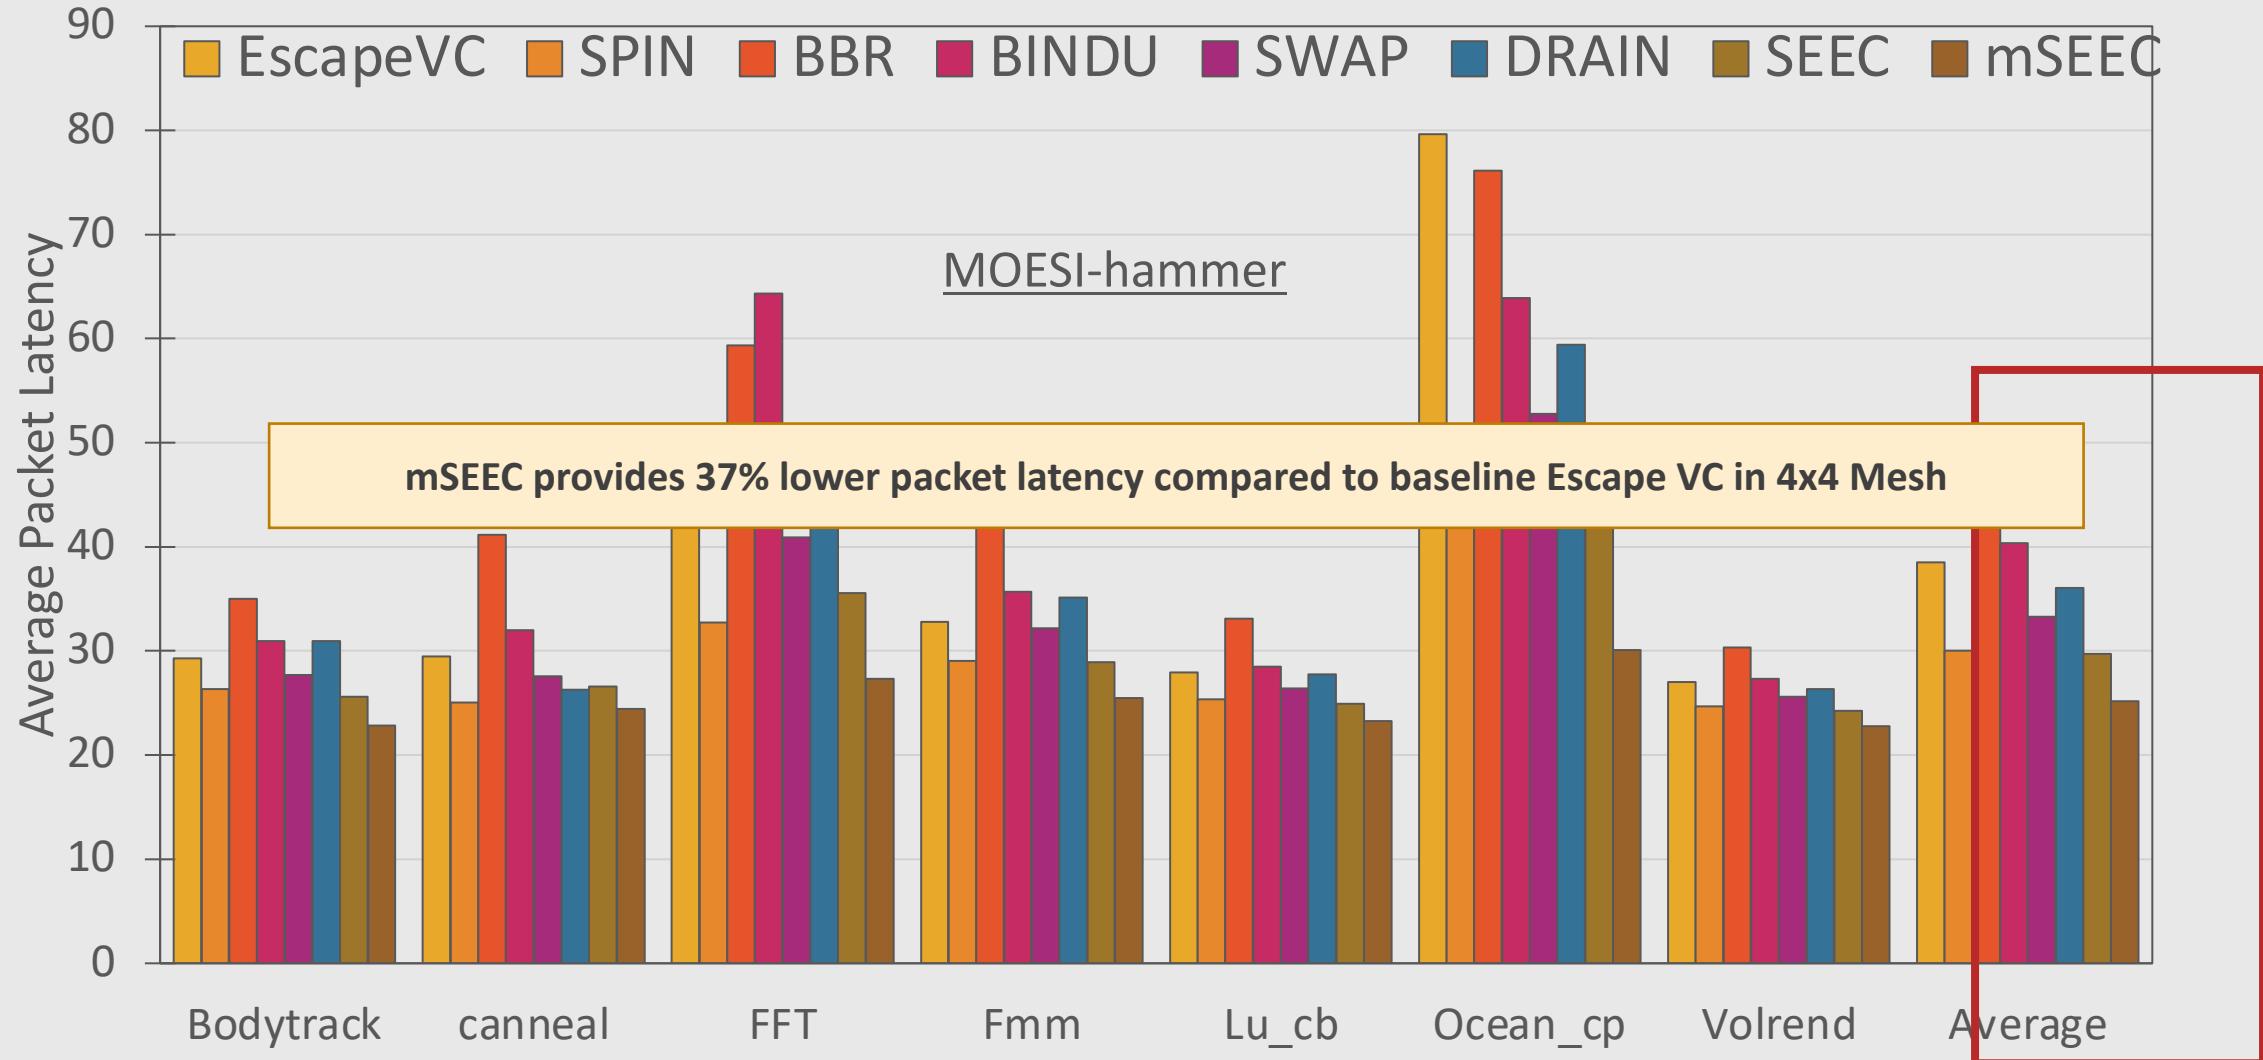
16% throughput improvement in 4x4 Mesh

Results: SEEC-%age of FF packets

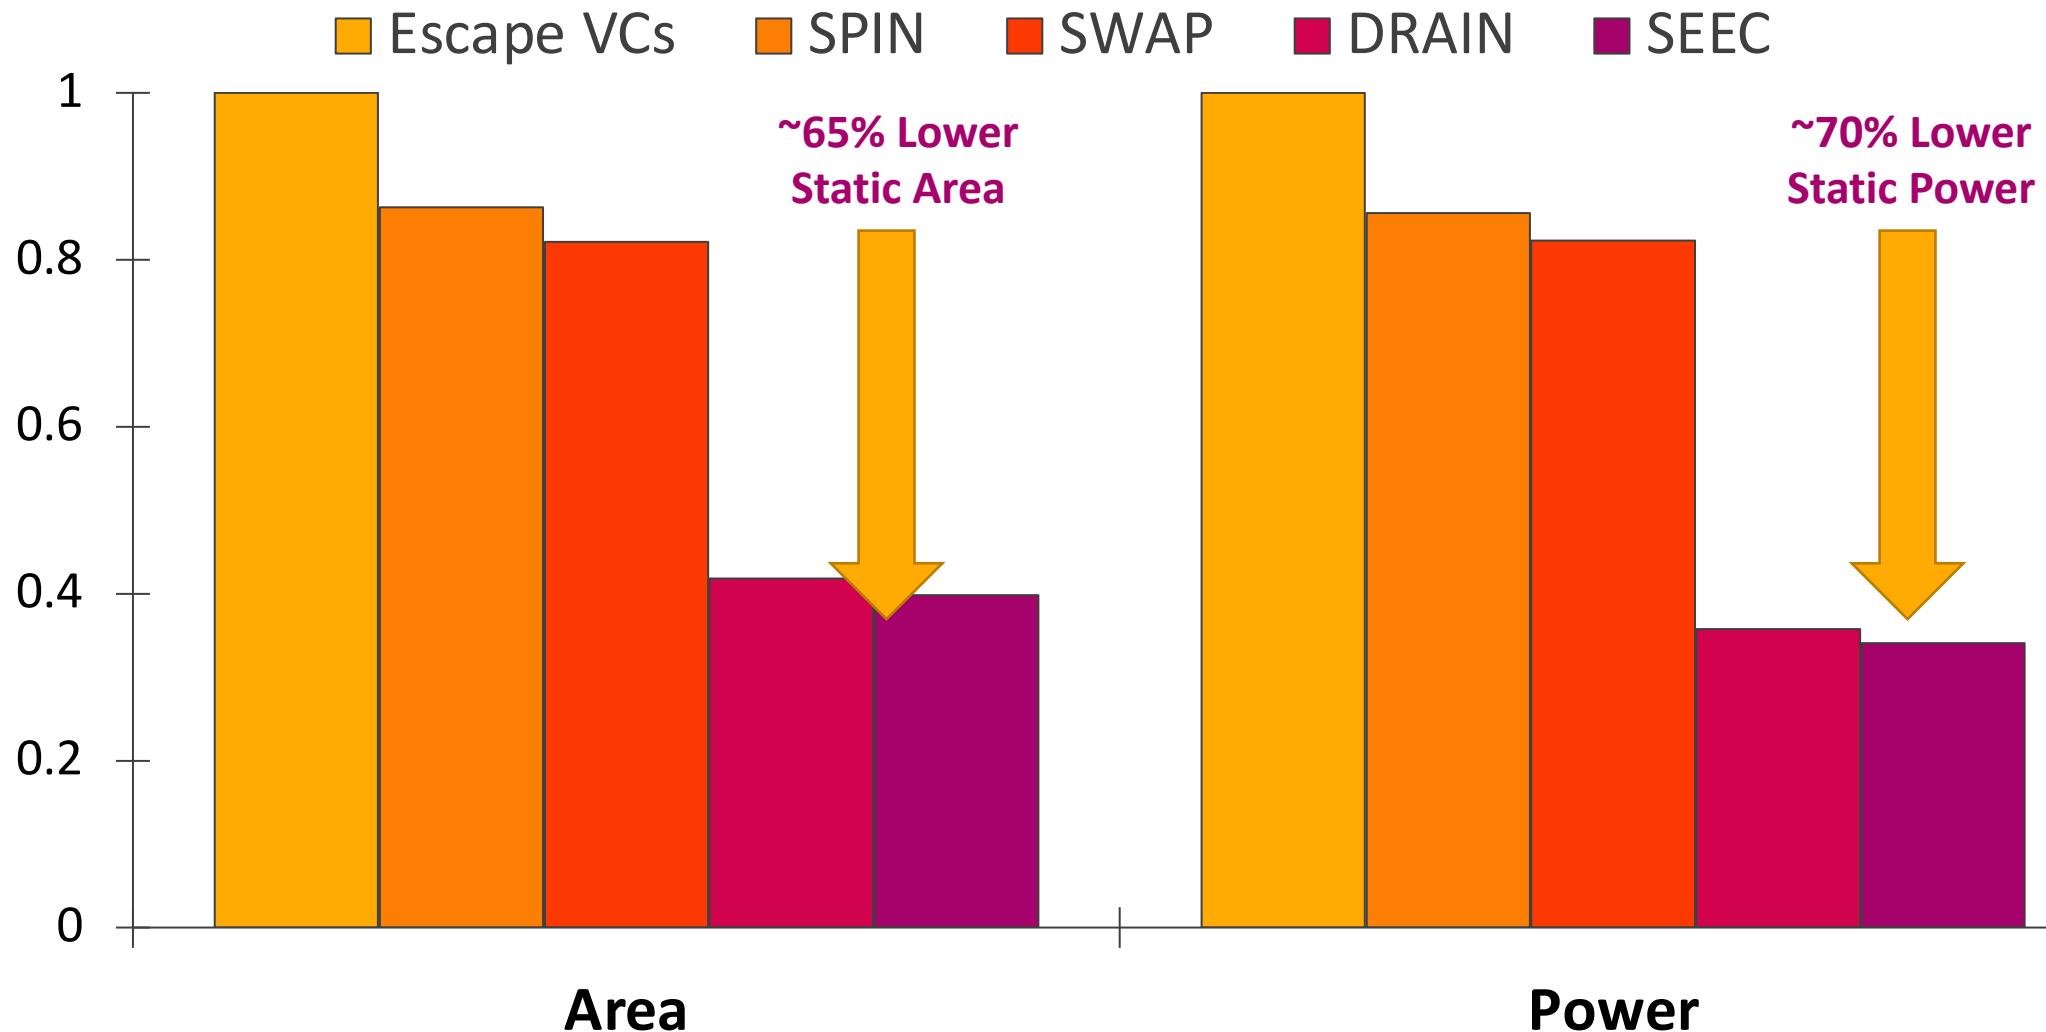


Post saturation heavy reliance on FF for packet delivery

Results: Full System Simulation



Normalized Area and Static Power



Outline

- **Background**
 - Networks-on-Chip
 - Deadlock-Freedom Techniques
- **SEEC**
 - Overview
 - Walkthrough
 - Features
 - Optimizations
- **Evaluations**
- **Conclusions**

Conclusion

- Deadlock is a correctness problem which every network must solve
- Current solutions either require complex deadlock detection or provision extra set of buffers or restrict turns
- SEEC provides a novel way of resolving deadlock by periodically moving packets for few hops along a ring constructed over the network.
 - No extra buffer needed
 - Free from complicated deadlock detection circuitry
 - Provides full path-diversity therefore high performant

Revisiting Qualitative Comparison

Codebase available here:
<https://github.com/noc-deadlock/seec>

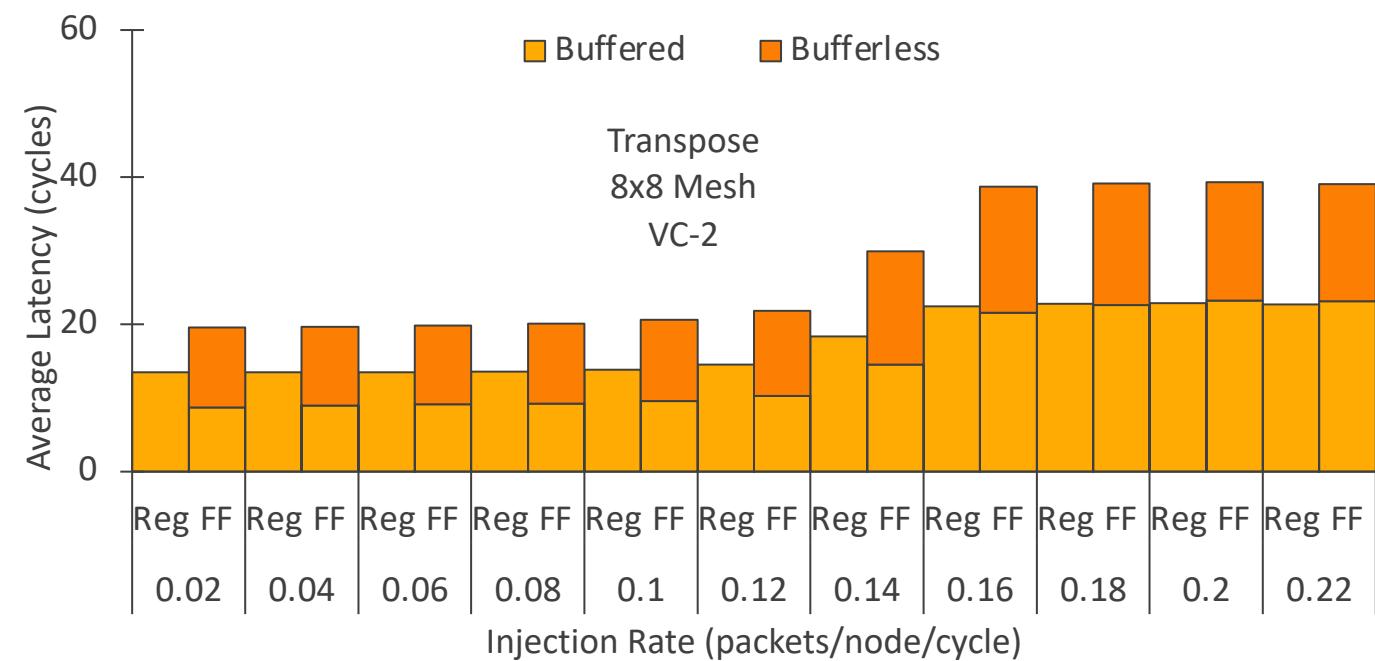
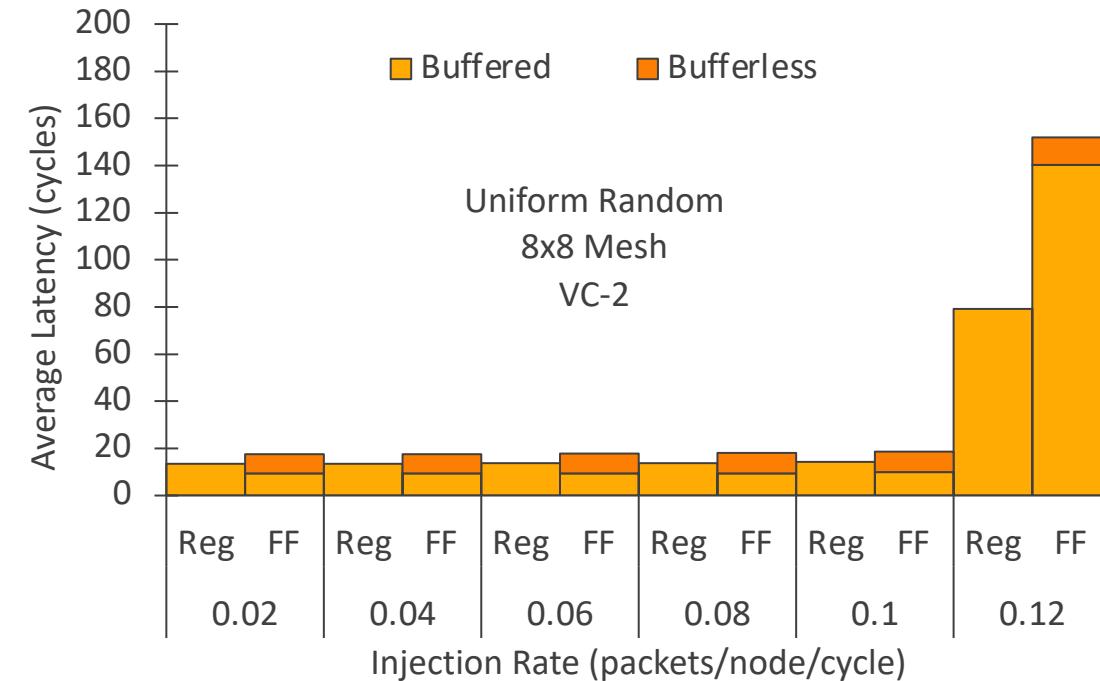
Technique	Type of Solution	High Throughput	Low Area Power	Low Hardware Complexity	Routing Deadlock-free	Protocol Deadlock-free
Turn Restriction	Proactive	✗	✓	✓	✓	✗
Escape VC	Proactive	✓	✗	✓	✓	✗
Virtual Network	Proactive	✓	✗	✓	✗	✓
Deadlock Recovery	Reactive	✓	✓	✗	✓	✓

Takeaways

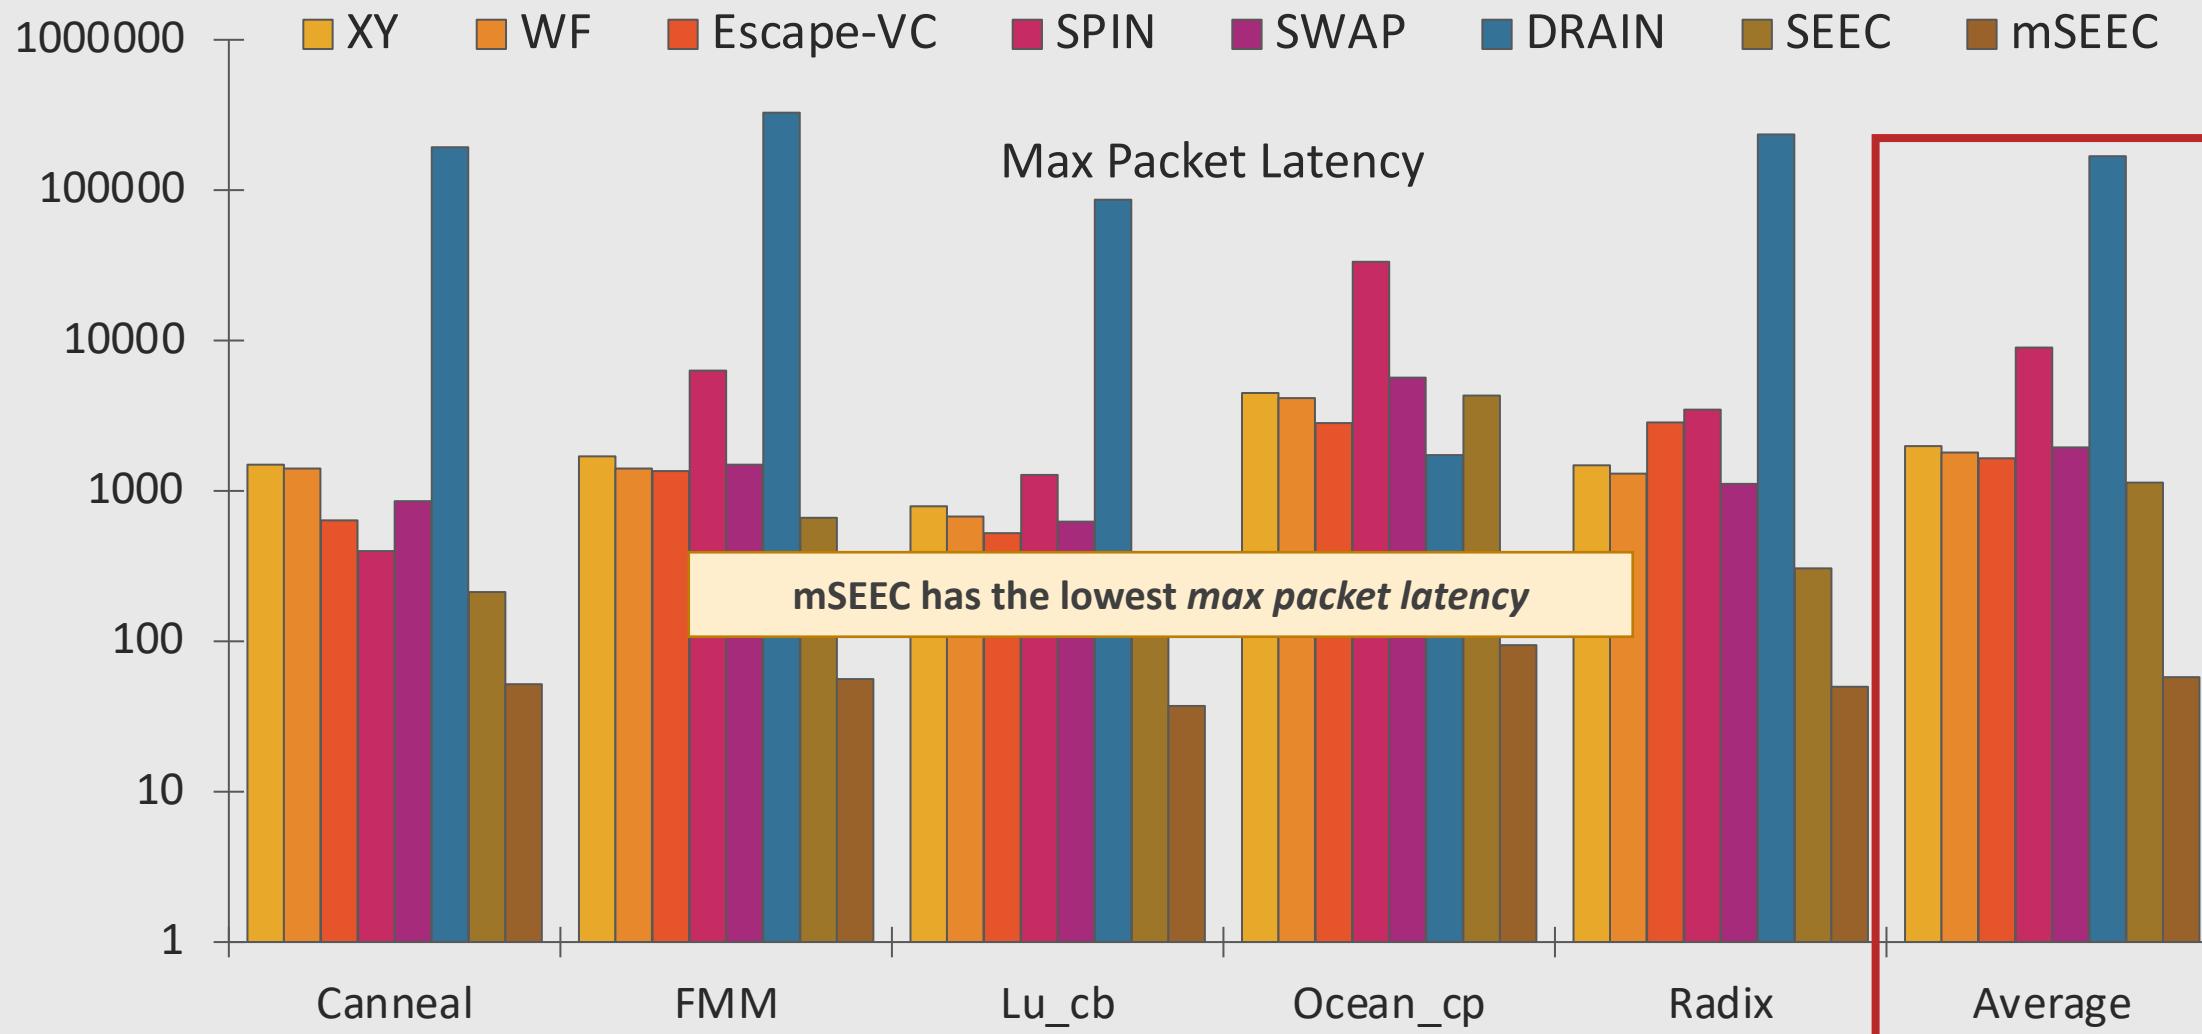
- SEEC obviously provides deadlock freedom
- SEEC is completely decentralized
- SEEC has low hardware overhead over baseline router

Thanks!

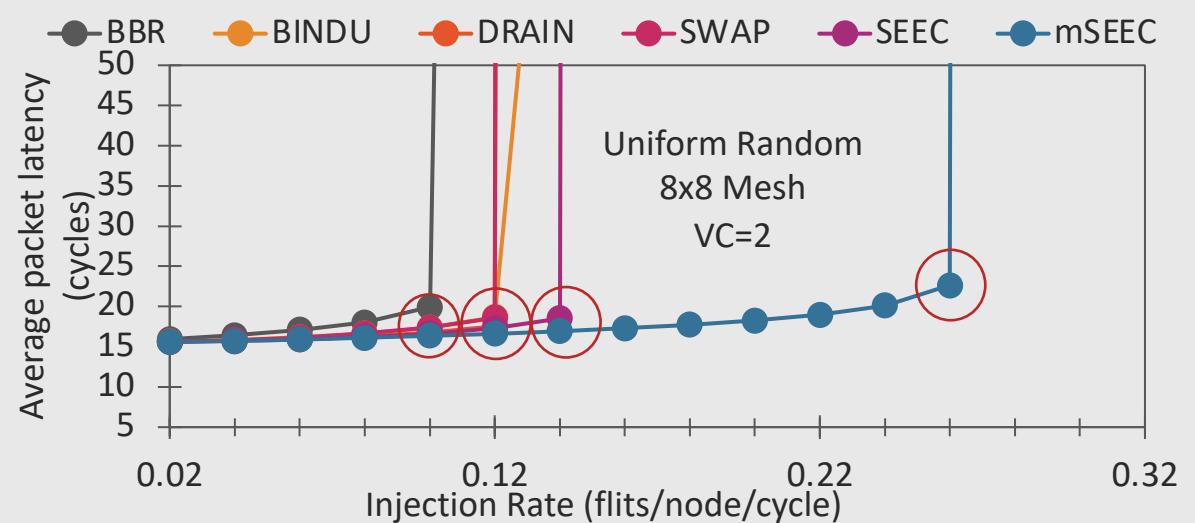
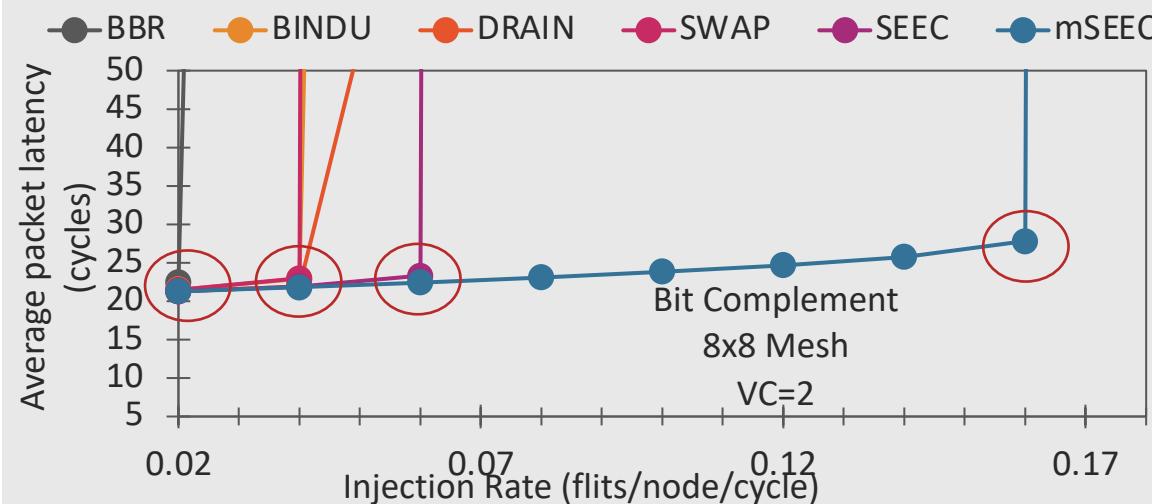
Results: SEEC-Network Latency breakdown



Result: Max Packet Latency



Results: Subactive Techniques



We observe following performance order (subactive techniques only) :
 BBR < BINDU ≈ DRAIN ≈ SWAP < SEEC < mSEEC

